



Asia's Largest

AI & Cloud

Conference 2024

15 - 16, November 2024

Chennai Trade Center, Chennai





It's me

Santosh Hegde

Senior Director- Couchbase Inc

About Me

17+ years of Database and Distributed Systems experience at IBM Software Labs and Visa Inc. Currently leading R&D at Couchbase.



Analytics beyond relational databases

Unlocking the power of schemaless data



Ever since there have been software applications that produced data...

the next logical step...

is to ask questions of that data, get answers, and improve the business

A vertical banner featuring a collage of abstract geometric shapes and patterns in blue, green, and white. The design includes stylized letters like 'n' and 'd', circles, squares, and floral motifs, all arranged in a modern, layered composition.

- ## This application evolution brings about new challenges for Analytics

Analytics Challenges | JSON Applications with RDBMS Analytics



1. Relational Rigidity

JSON's developer flexibility does not align with traditional relational analytic structures



2. Multi-source Data

Silos of data must be consolidated from unstructured and structured origin sources



3. Slow Analytic Turnaround

Days or weeks to get changes for new metrics, queries and reports from BI / Data Warehouse team



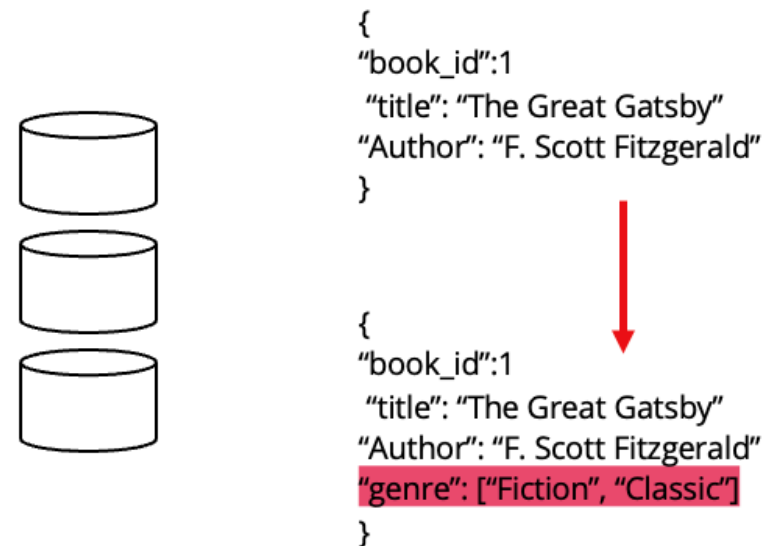
4. No Actionability

Difficult to incorporate newly learned information based on calculated metrics to improve application functionality and AI conversations

Capella Columnar is designed to overcome these challenges

Analytics Challenges | Relational Rigidity

JSON Transactional Data



Rigid Analytic Warehouse

Books		
Book_ID	Title	Author
1	The Great Gatsby	F. Scott Fitzgerald

Genres	
ID	Type
1	Fiction
2	Classic

BookToGenre	
Book ID	GenreID
1	1
1	2

JSON not intended for Analytics

- Unsupported structures:
 - Subdocuments
 - Arrays, Geo coordinates

Built for OLTP applications

Imports must match fixed schema

- Central fact table with star and snowflake-style dimensions
- Data integrity-driven

Built for OLAP analytics

Analytics Challenges | Multi Source Data


JSON & other data sources



Analytics have multiple sources

- **JSON (nested and heterogeneous)**
- Flat files are common
- RDBMS data are common
- Application data is real-time data

Multiple Ingestion Pipelines & File Imports



ETL & data pipelines hate change

- New fields in JSON documents
- Null and missing values
- Fields of different data types
- **All break the pipeline**

Rigid Data Warehouse

Books		
Book_ID	Title	Author
1	The Great Gatsby	F. Scott Fitzgerald

Genres	
ID	Type
1	Fiction
2	Classic



BookToGenre	
Book ID	GenreID
1	1
1	2

This leads to extra work

- This can require additional fields and tables to be created in the warehouse
- Breakage is common

Analytics Challenges | Slow Analytic Turnaround

**Application
team**



**BI/DW
team**



-
- Requests for report and analytics changes must go through another team bureaucracy (e.g., Data Engineers, BI Analysts, BI Consumers)

-
- Analytics team handles requests from many teams across the organization

-
- Need to balance requests leads to time delays
 - Often, they don't know the data as well as the application team

Analytics Challenges | No Actionability



- Analytic stores are not designed to writeback to transactional databases

- **Reverse ETL programs or custom code** needed to extract metrics data to the transactional database

- Often costly and time consuming to set up and maintain



Capella Columnar | Next Generation JSON Analytics

Capella Columnar

NoSQL

- Schema Flexibility
- High Availability
- Scalability
- Developer Friendliness



Analytical


- Columnar Storage
- Cost Based Optimizer
- Massive Parallel Processing
- Storage/Compute Separation

Capella Columnar | Next Generation JSON Analytics



Zero ETL for JSON

- JSON automatically processed
- Schema on Read
- Changes handled
- Transformations happen automatically



Multi-source Ingestion

- Capella & Couchbase Server
- MongoDB
- MySQL & PostgreSQL
- Files: CSV, TSV, Parquet, Avro, etc.
- Batch and streaming



Conversational Analytics

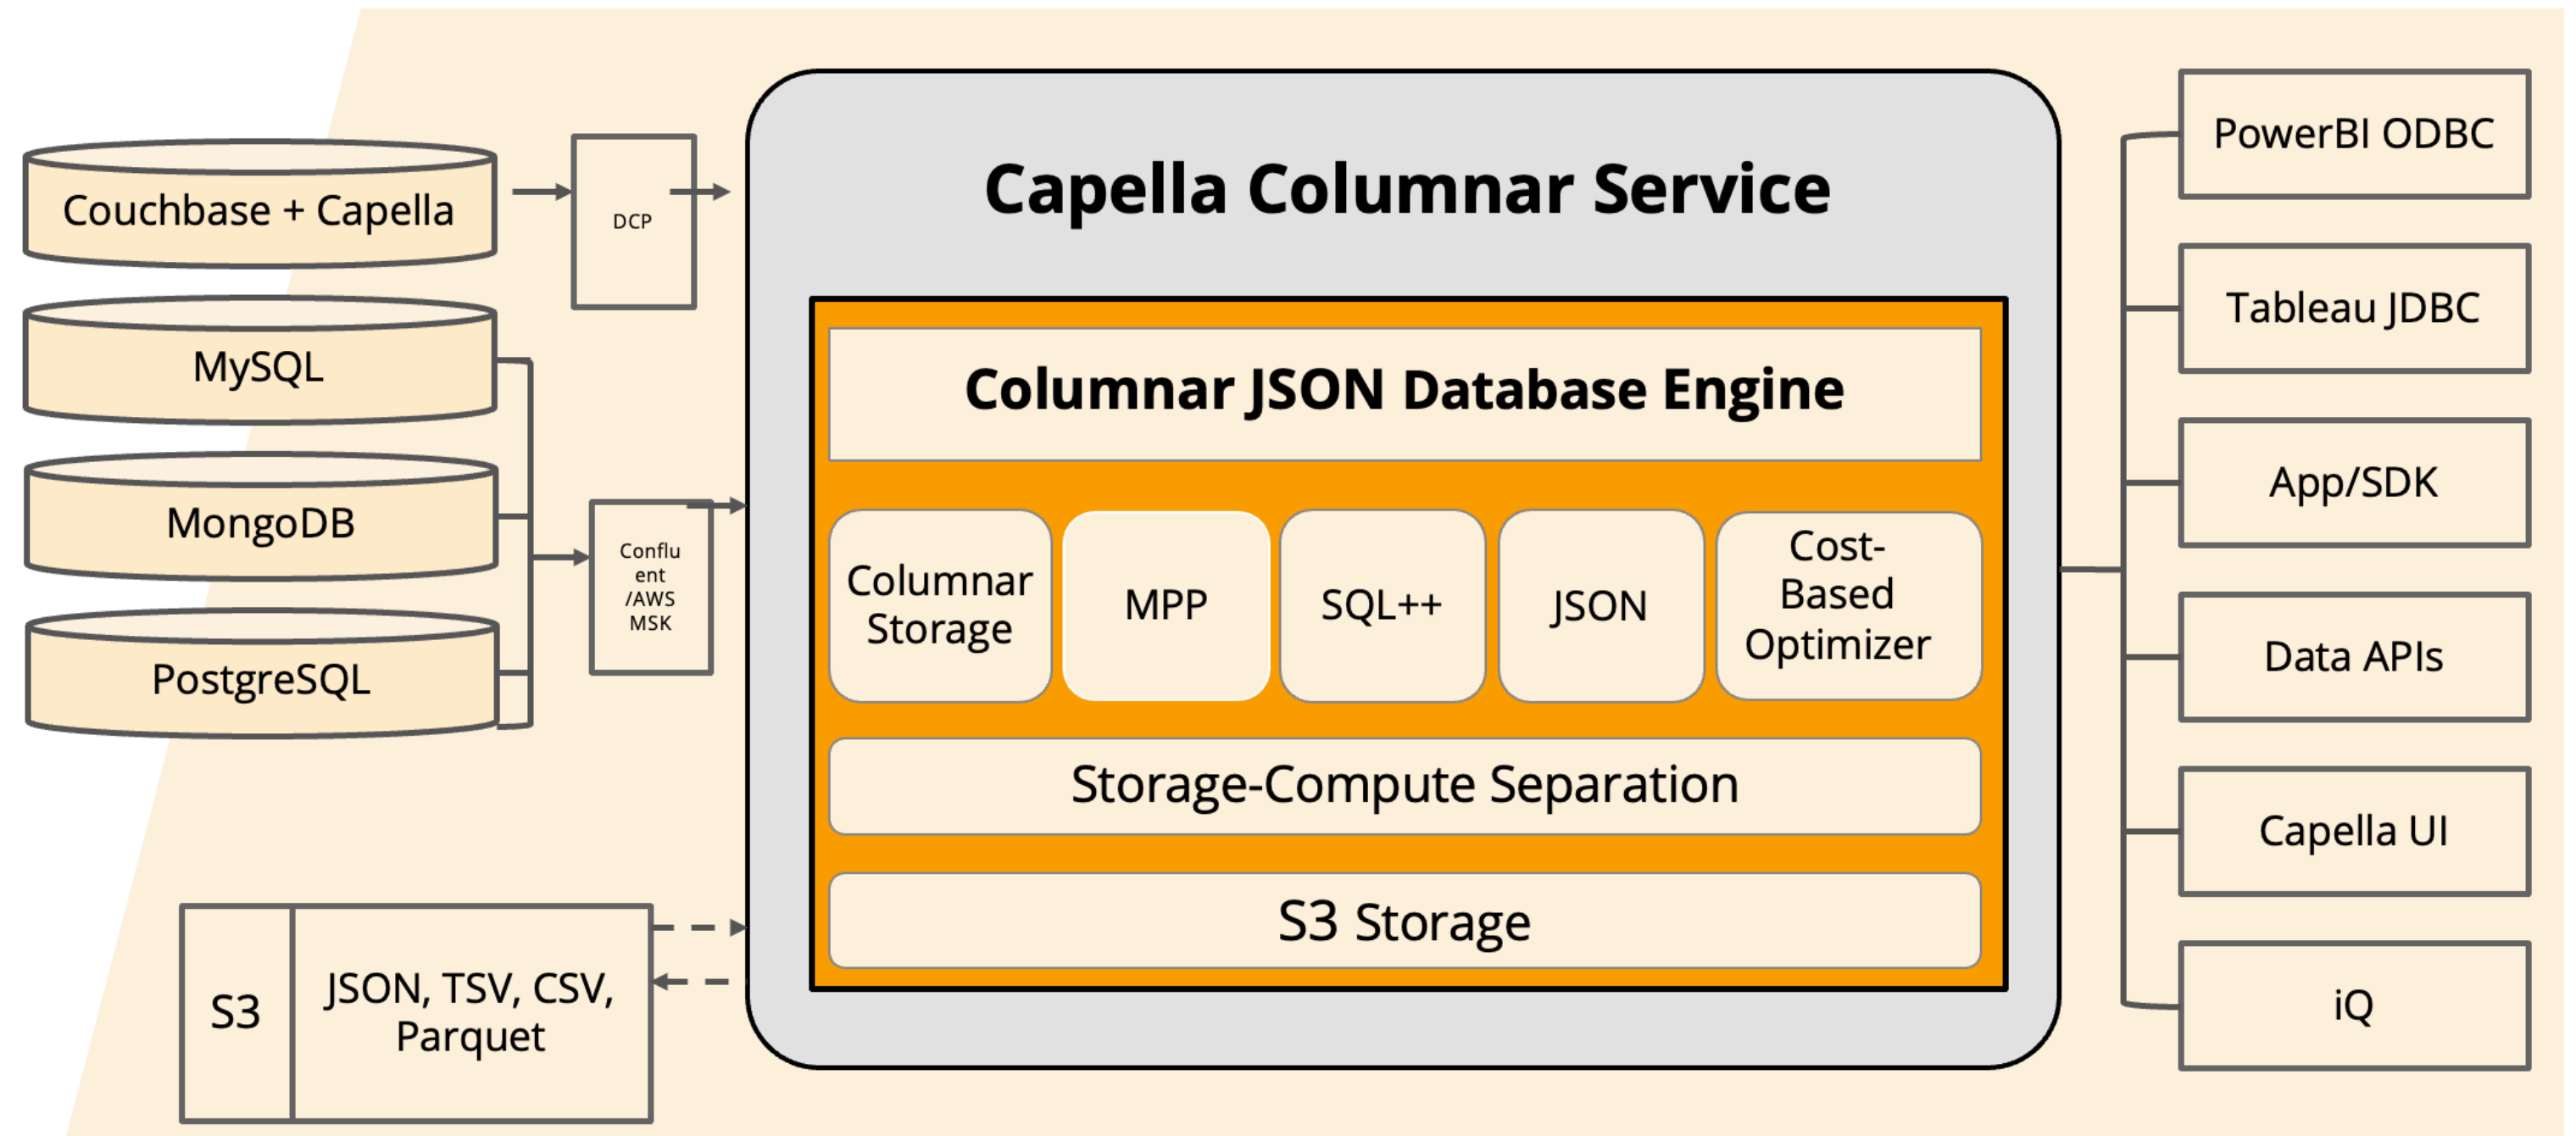
- Developer-controlled analytics using iQ
- Capella with Columnar for operational and analytical applications



Operational Writeback

- Drive action from calculated metrics
- Feed AI prompts
- Build trustworthy AI-enabled features

Capella Columnar | Service Architecture



Capella Columnar | What does it enable

More Personalized User Experiences

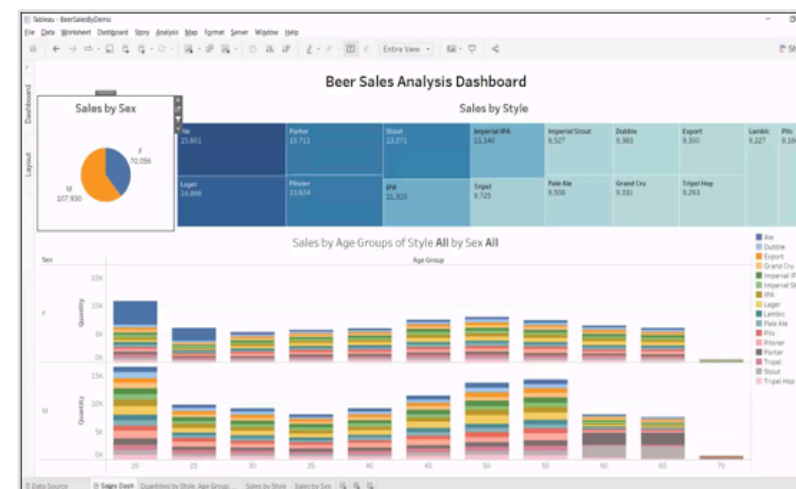
- Contextualized user experiences incorporating AI
- Driving custom offers



"You've hit your saving goal, book the trip!"

Operational Analytics

- Operational visibility
- Smoother DevOps
- Faster anomaly detection




Customer-Facing Metrics

- Embedded analytic information
- For B2B or user-facing apps





Questions ??





THANK YOU!

Santosh Hegde

Email: santosh.hegde@couchbase.com

LinkedIn: [santosh-hegde-cb](https://www.linkedin.com/in/santosh-hegde-cb)

