



Architecting the future: Patterns to build generative AI application

Praveen Jayakumar

Head of AI/ML Solution Architecture
AWS India

LLM customisation patterns

Prompt engineering (In-context learning)

Retrieval Augmented Generation (RAG)

Fine-tuning


Training your own LLM



Prompt engineering types


Zero shot prompts

- Direct request with sufficient context



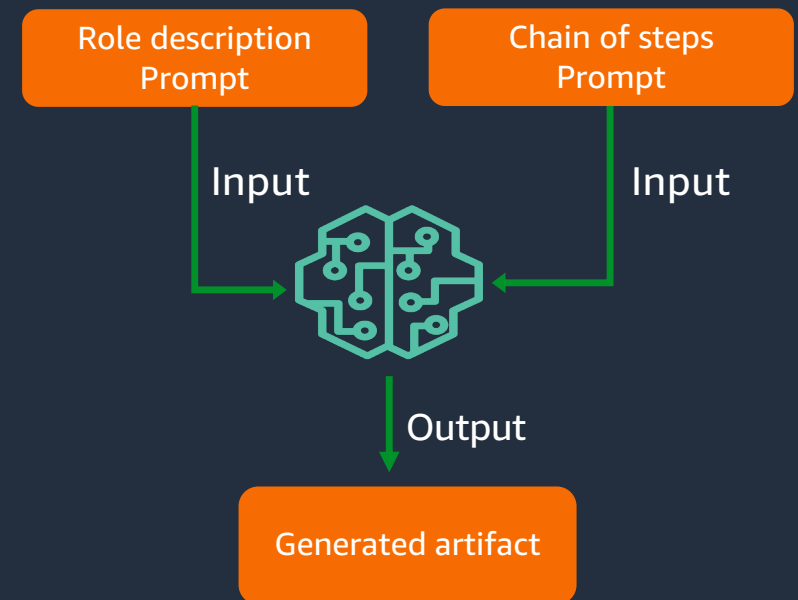
One shot or few shot prompts

- Provide one or more examples with a request



Role or Chain of Thought prompts

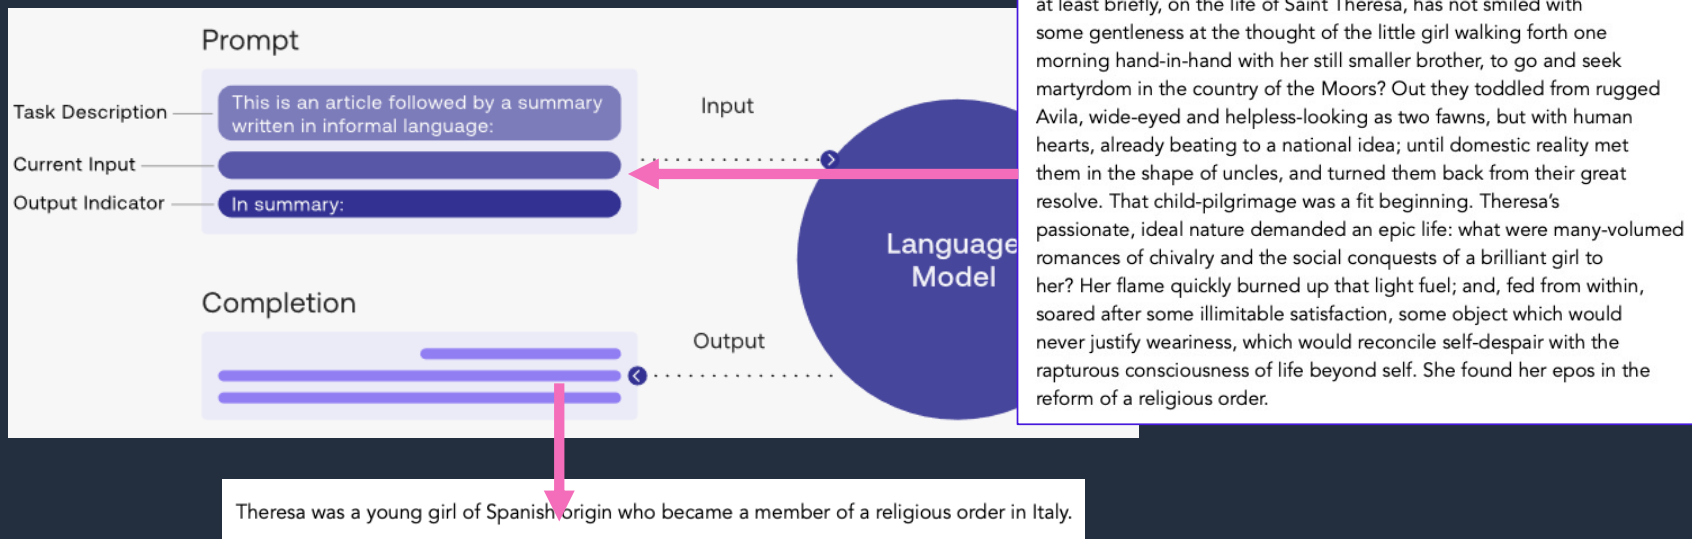
- Provide the model with a role or persona for the task
- Provide a chain of steps for the model to follow



Zero shot prompt: Prompting by instruction

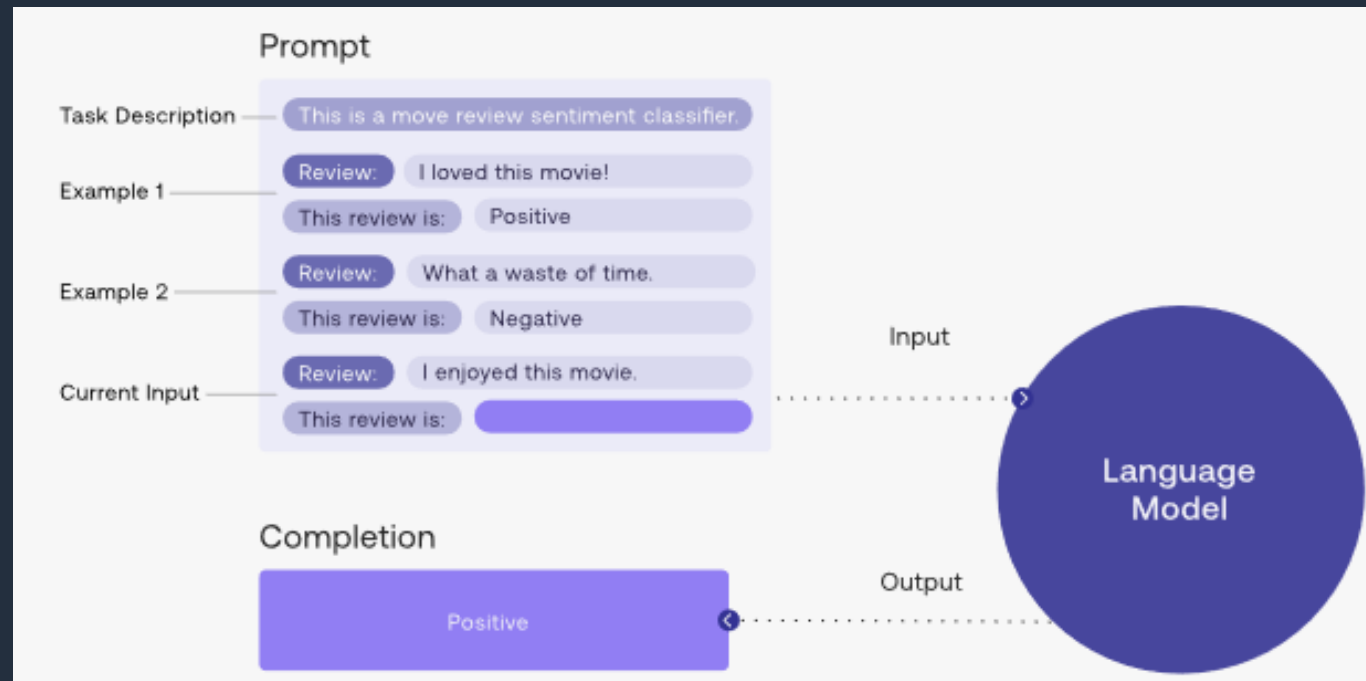
Zero-shot prompting:

- Using LLM out of the box
- allows language models to perform tasks for which they have not been explicitly trained on



Prompt Engineering: N shots example

- A **few-shot** prompt normally includes **n examples** of (problem, solution) pairs known as "**shots**".
- Help to guide **model performance**.



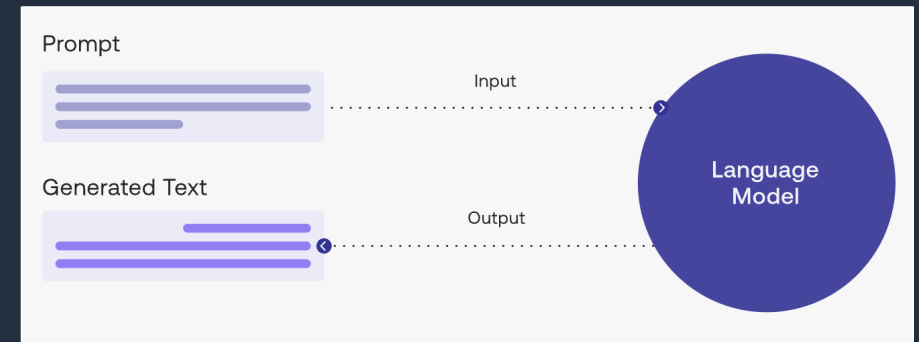
Prompt Engineering: Chain of thought prompting example

- Improves *reasoning* abilities in foundation models
- Addresses *multi-step problem-solving* challenges in arithmetic and *commonsense reasoning* task
- Generates intermediate reasoning steps, mimicking *human train of thought*, before providing the final answer.

Standard Prompting	Chain-of-Thought Prompting
<p>Model Input</p> <p>Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?</p> <p>A: The answer is 11.</p> <p>Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?</p>	<p>Model Input</p> <p>Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?</p> <p>A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.</p> <p>Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?</p>
<p>Model Output</p> <p>A: The answer is 27. ❌</p>	<p>Model Output</p> <p>A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✅</p>

Prompting limitation

- Poor memory
- Limited context
- Accessing external knowledge sources to complete tasks.



What is Retrieval Augmented Generation?



Retrieval

Fetches the relevant content from the external knowledge base or data sources based on a user query



Augmentation

Adding the retrieved relevant context to the user prompt, which goes as an input to the foundation model

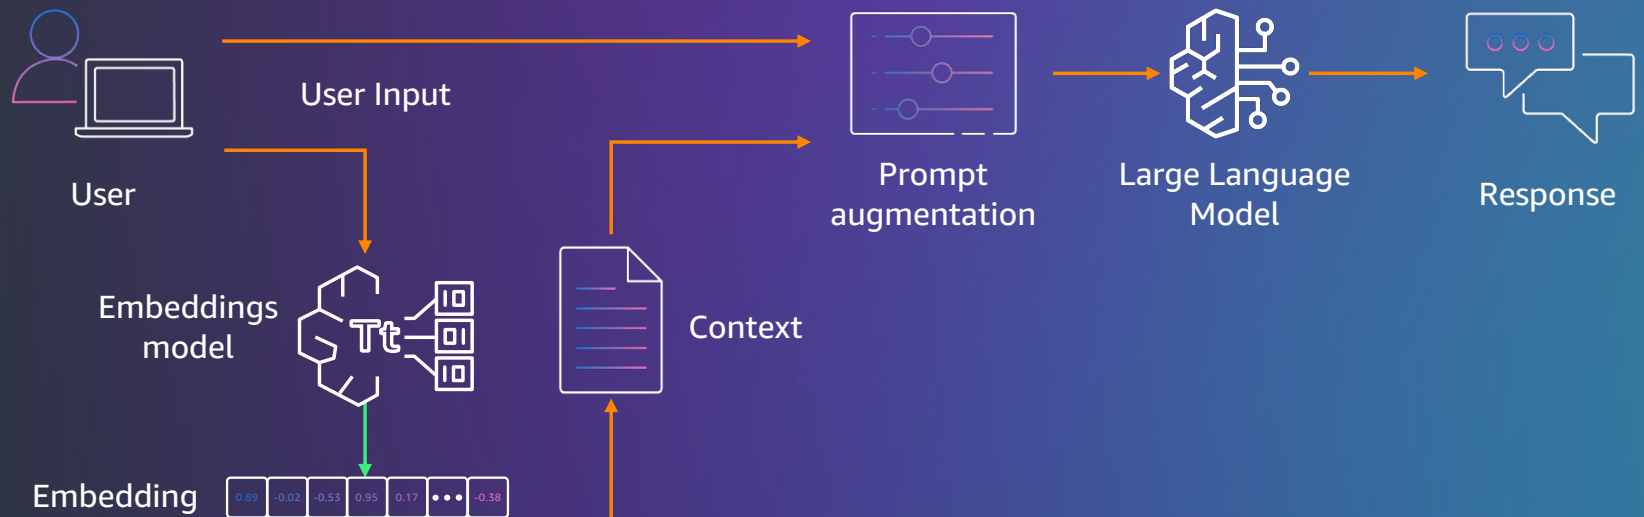


Generation


Response from the foundation model based on the augmented prompt.

RAG in Action

Text Generation Workflow



Data Ingestion Workflow



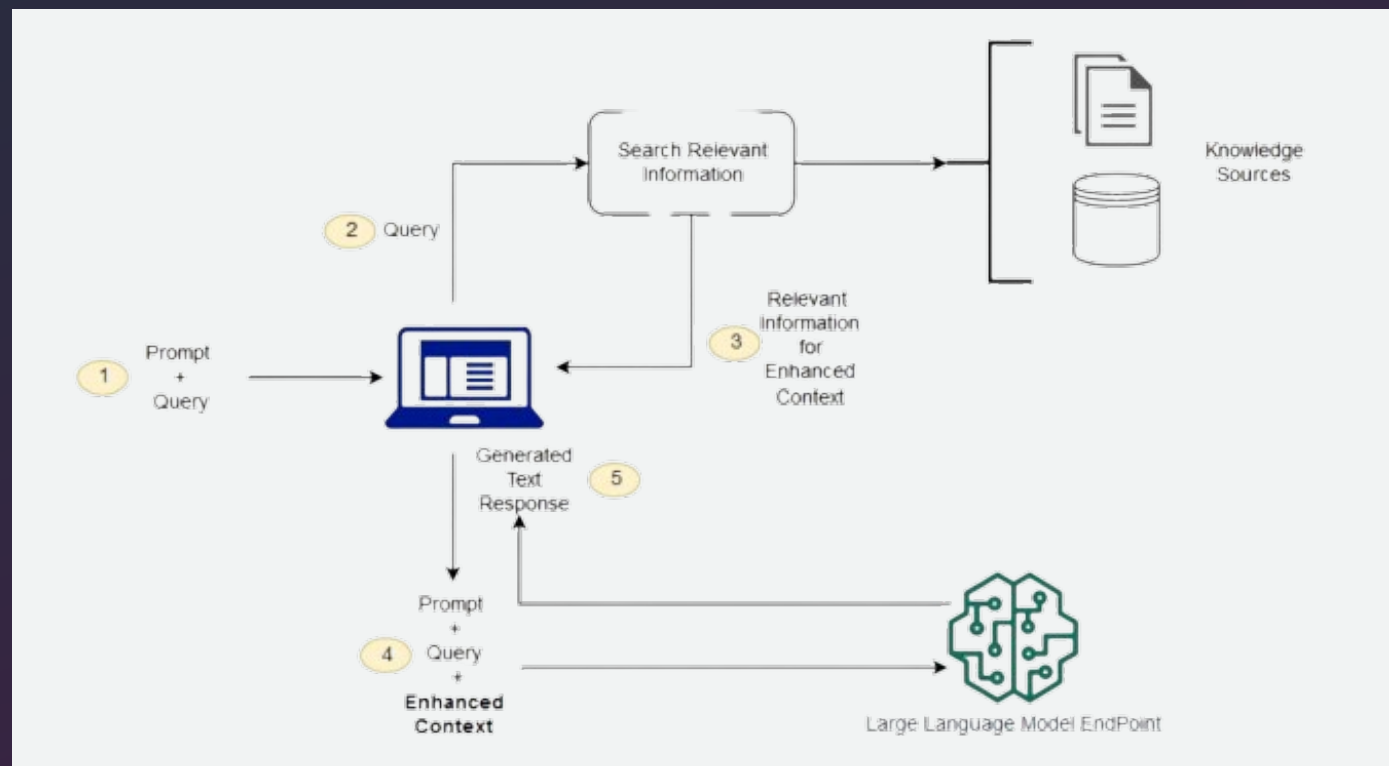
RAG Patterns



© 2024, Amazon Web Services, Inc. or its affiliates. All rights reserved.

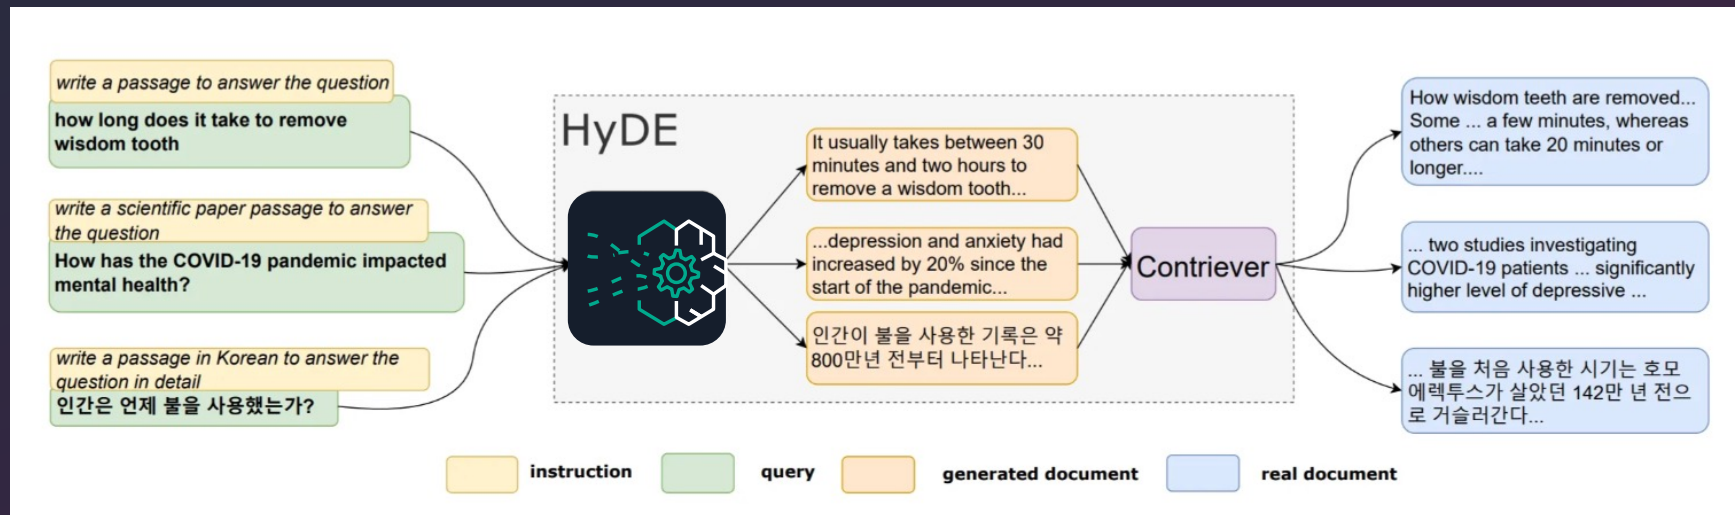
Simple RAG

1. Query
2. Retrieve Docs
3. Prompt + [Docs]
4. Generate Response



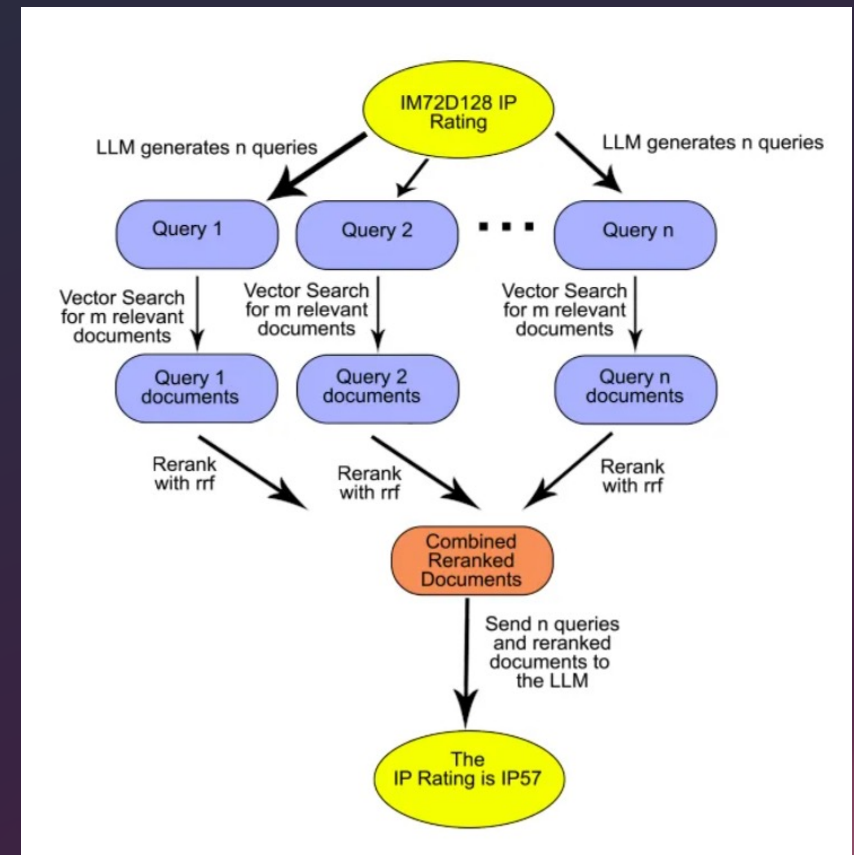
HyDE: Hypothetical Document Embedding

1. Query
2. Generate Hypothetical Answer
3. Embed Hypothetical Answer
4. Retrieve Docs with Embedding
5. [Docs]
6. Generate Response



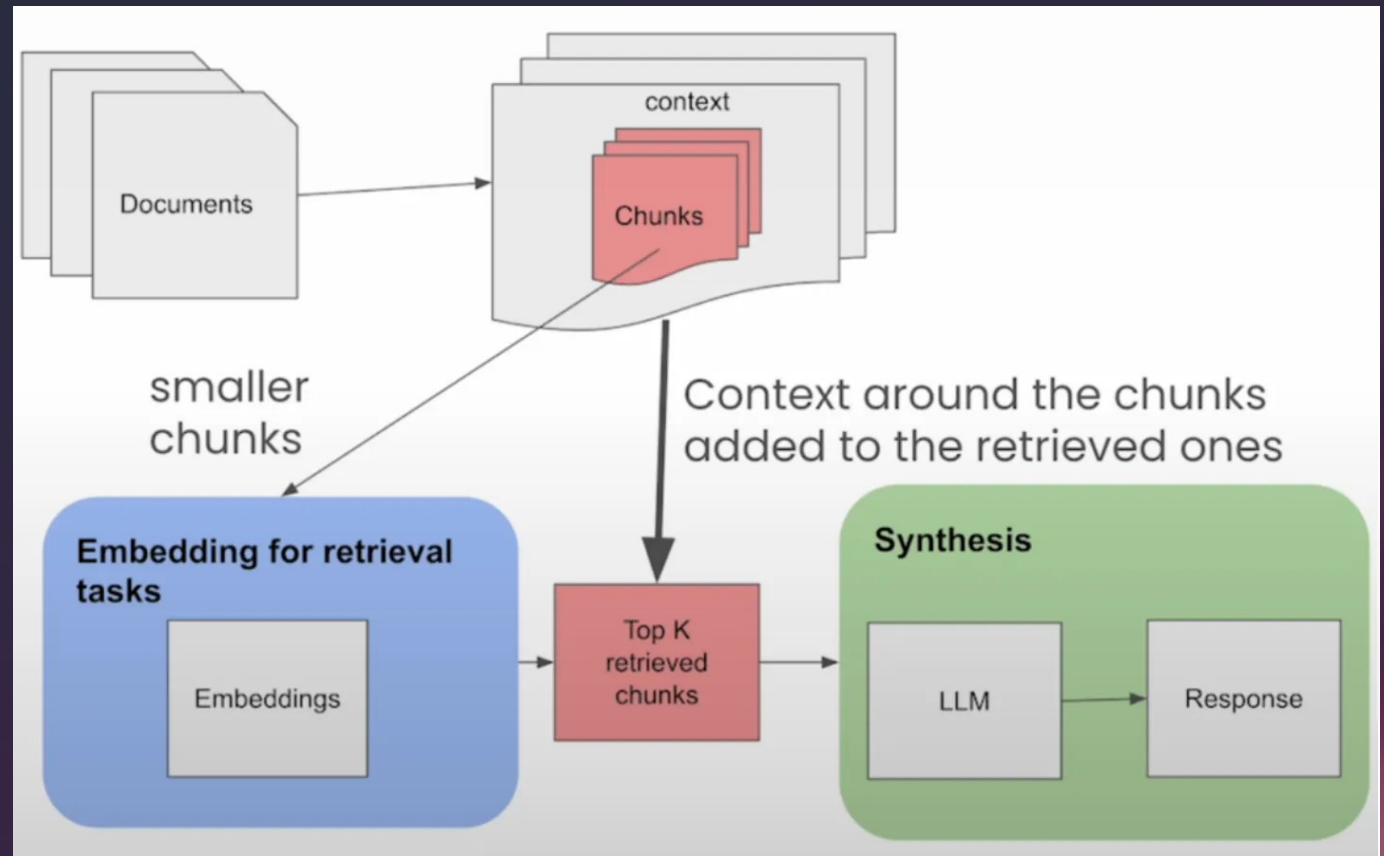
Multi Query RAG

1. Query
2. Generate Multiple Queries
3. [Retrieve Docs(query 1), Retrieve Docs(query 2), ...]
4. Rerank Docs – Reciprocal rank fusion
5. [Reranked Docs]
6. Generate Response



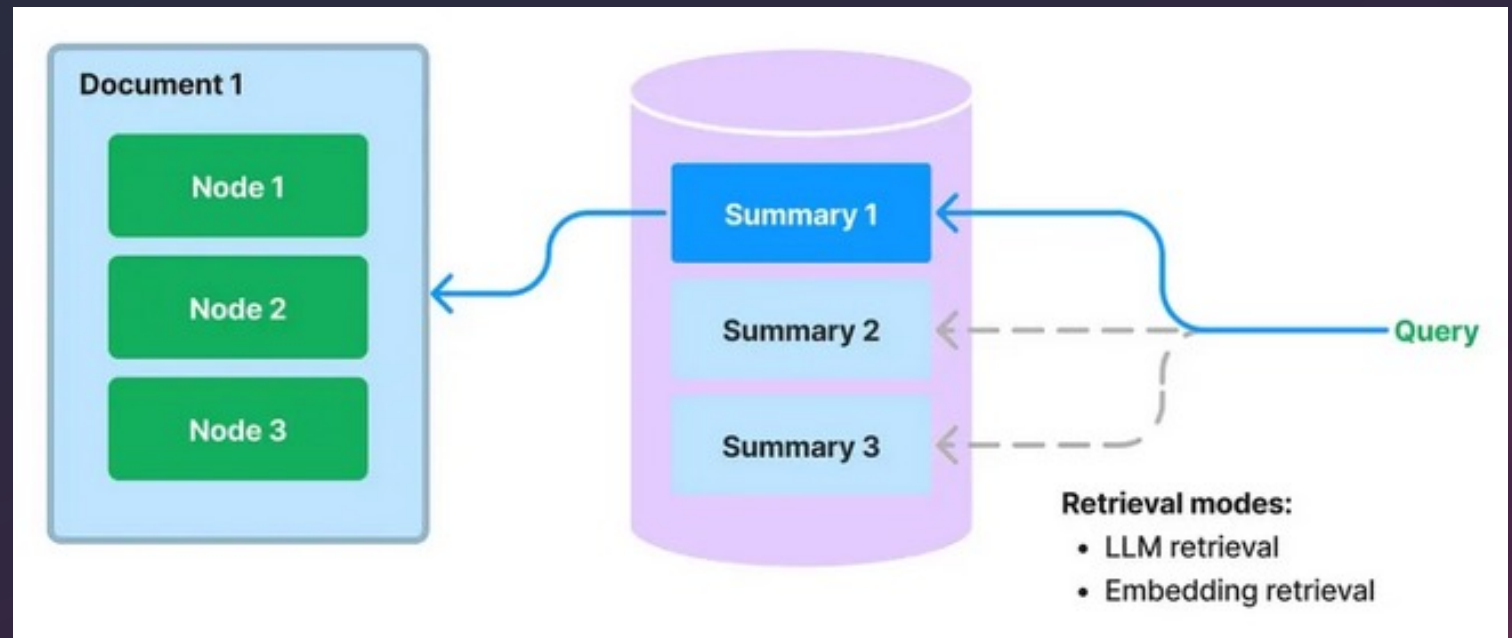
Sentence Window Retrieval

1. Query
2. Retrieve Sentence Windows
3. [Sentence Windows]
4. Add Context around Sentence Windows
5. Generate Response



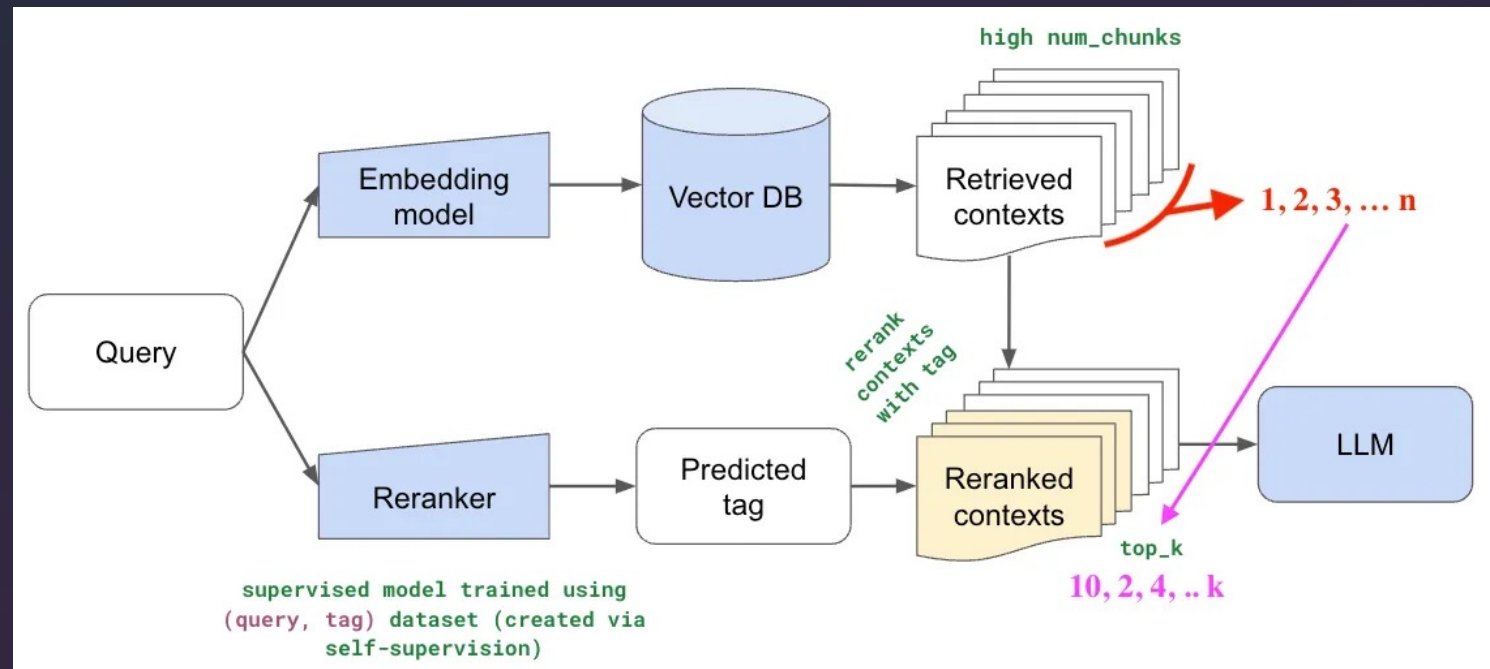
Document Summary Index

1. Query
2. Retrieve Document Summaries
3. [Full Documents]
4. Generate Response



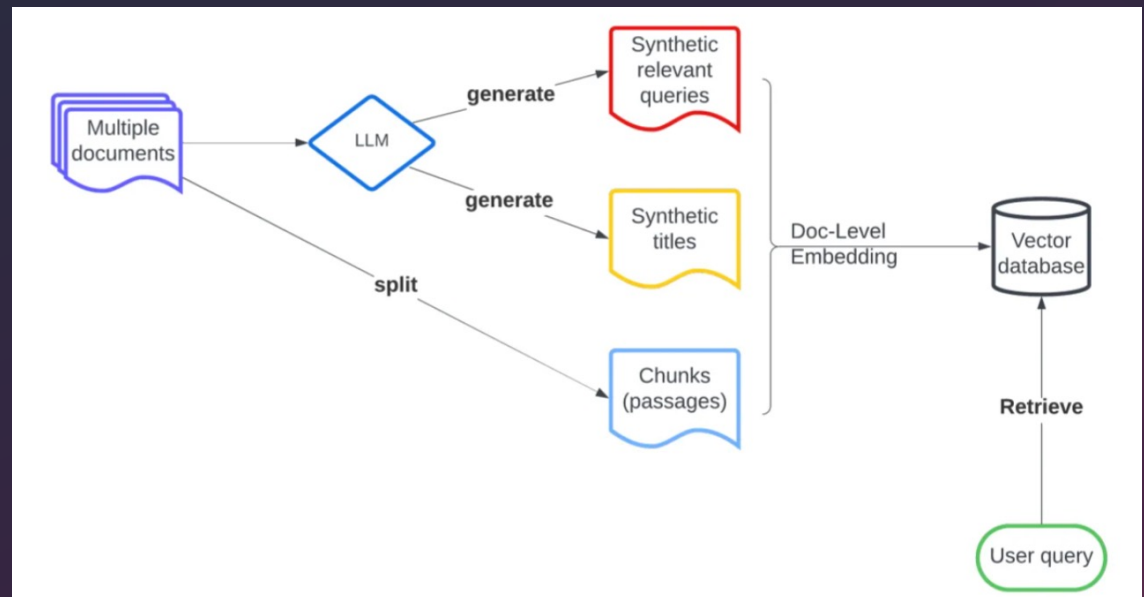
Reranker (MMR, Cohere, LLM)

1. Query
2. Retrieve Docs
3. Rerank Docs
4. [Reranked Docs]
5. Generate Response




LLM-Augmented Retrieval

1. Documents
2. Split
3. [Chunks]
4. LLM
5. Generate [Synthetic Titles]
6. Generate [Relevant Queries]
7. Combine ([Chunks, Synthetic Titles, Relevant Queries])
8. Doc-Level Embedding
9. Vector Database
10. Retrieve ([Relevant Docs])
11. User Query



FLARE - Forward Looking Active RAG

1. Documents
2. Split
3. [Chunks]
4. LLM
5. Generate [Temporary next sentences]
6. Combine Search results
7. Retrieve [Relevant Docs]
8. User Query

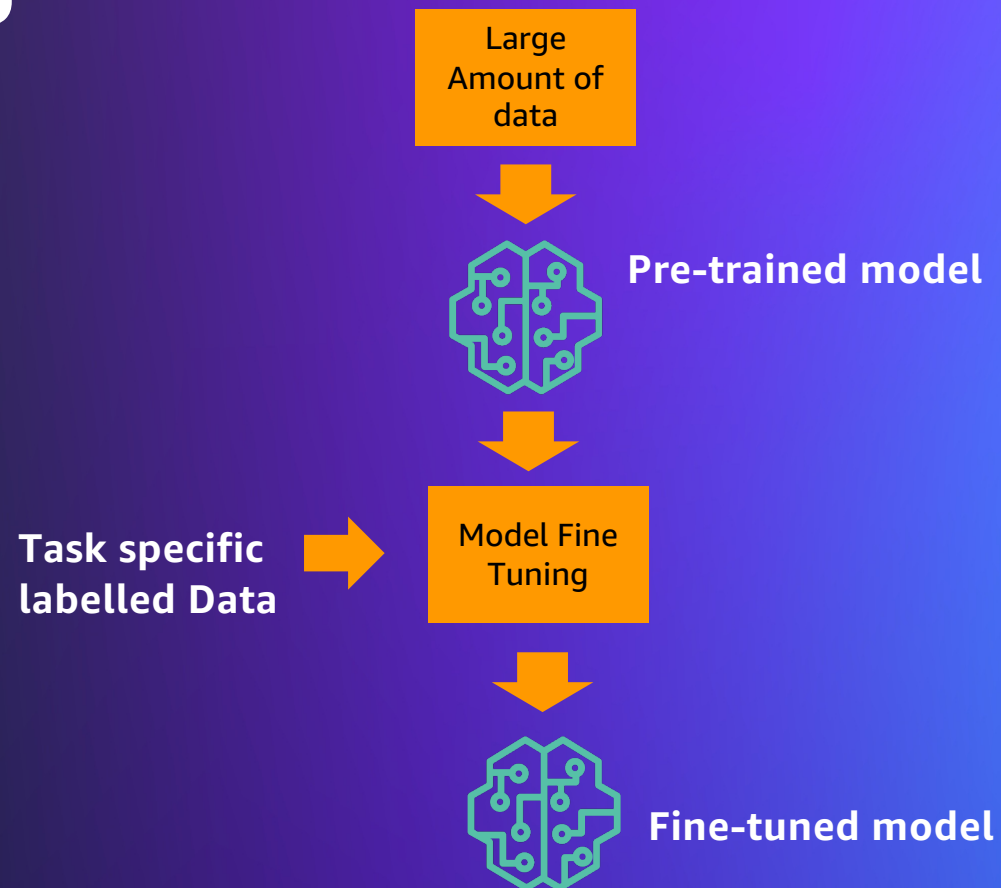


Key challenges with Retrieval Augmented Generation (RAG)

- Managing multiple data sources
- Creating embeddings for large volume of data
- Incremental updates to vector database
- Increased Complexity => adding retriever component to generation model
- Limited Creativity => constrained by the retrieved information
- Context window limited by LLMs



Fine-tuning



Fine-tuning

Traditional fine-tuning

- Expensive and resource intensive for LLMs scaling to multibillion parameters.
- Overfitting and Knowledge degradation: also known as "catastrophic forgetting"

PEFT fine-tuning

- Training only a subset of parameters => Faster, cheaper, less computational power
- Set of newly added parameters or select some existing model parameters.
- Most of existing weights do not change :
 - Reduces overfitting and catastrophic forgetting
- Fast improving techniques such as LORA and others











Limitation of fine-tuning

- Deep knowledge
- Time consuming and resources
- Need enough data to fine-tune
- Data could be constantly changing
- Some LLM model are not available for fine-tuning

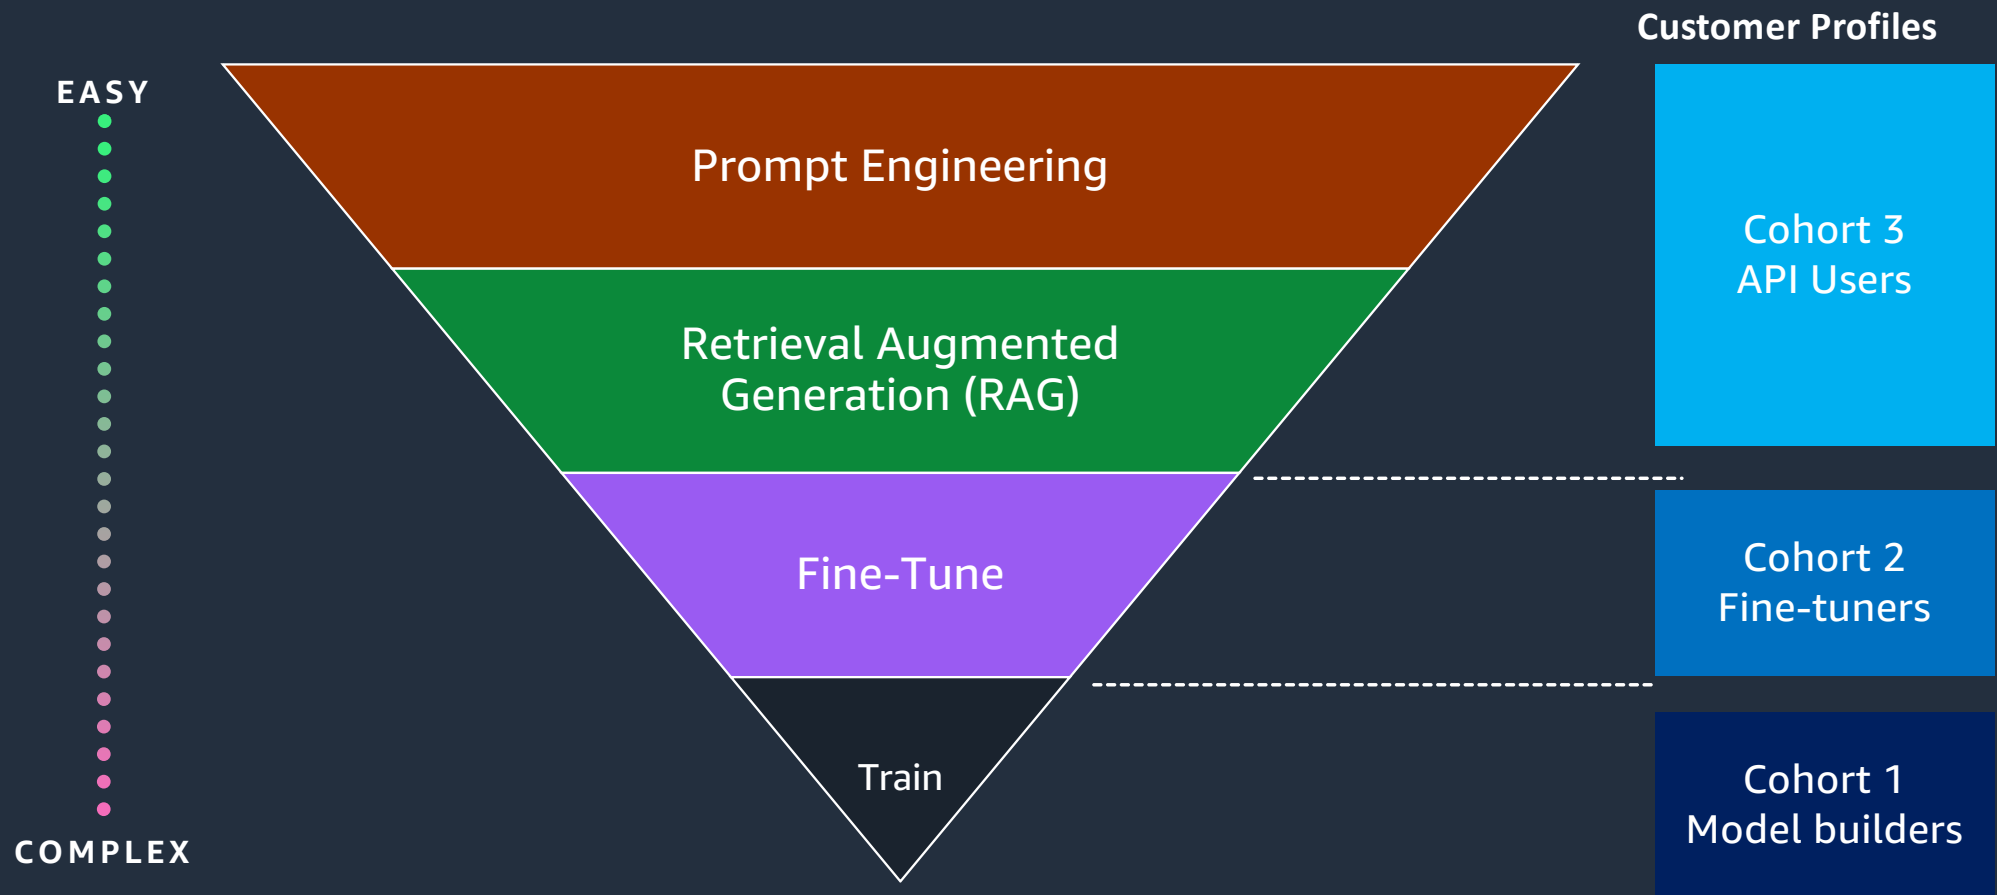


When should customers fine-tune

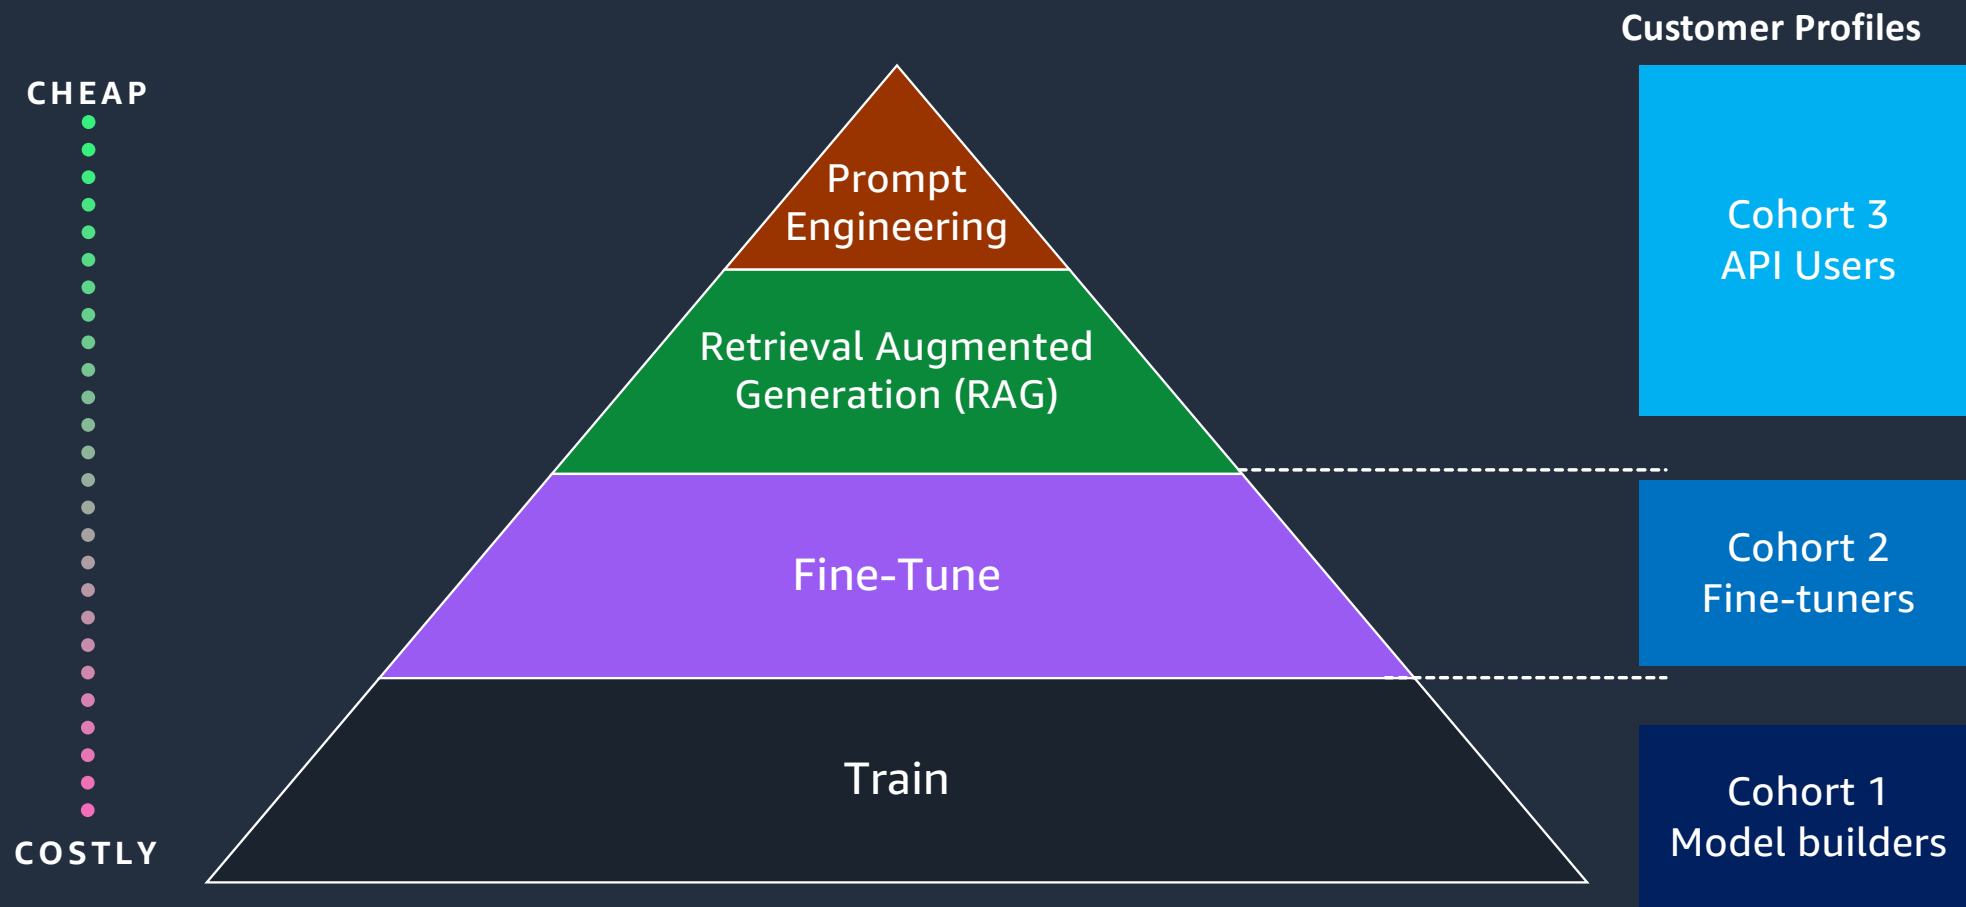
Problem	Example	Likelihood of success with fine-tuning	Likelihood of success with prompting (+ RAG)
Make the model follow a specific format or tone.	Use this specific JSON schema or talk like my customer service reps.	 VERY HIGH	 HIGH
Teach the model a new skill.	Learn how to call APIs, fill out proprietary documents, or classify customer support tickets.	 HIGH	 MEDIUM
Teach the model a new skill, and hope it learns similar skills.	Teach the model to summarize contract documents and hope it learns how to write better contract documents.	 LOW	 MEDIUM
Teach the model new knowledge and expect it to use that knowledge for general tasks.	Learn my company's acronyms or know more facts music.	 VERY LOW	 MEDIUM



LLM customisation comparison: Skills required




LLM customisation comparison: Cost



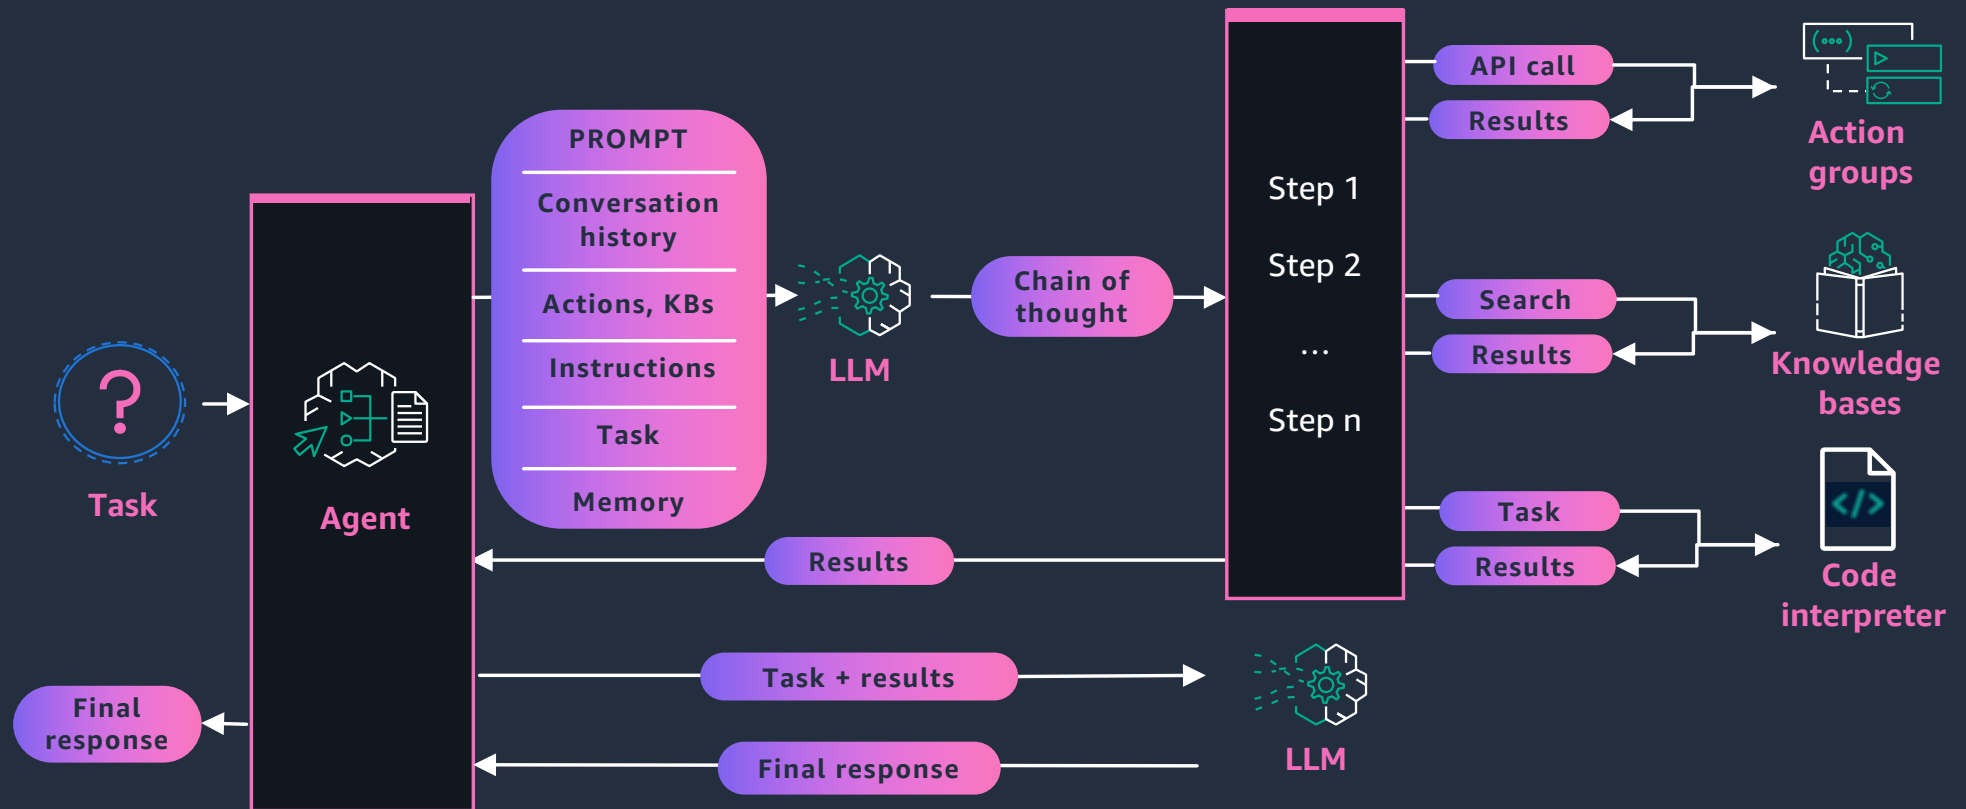
I think **AI agentic workflows** will drive massive AI progress this year — perhaps even more than the next generation of foundation models. This is an important trend, and I urge everyone who works in AI to pay attention to it. –
– **Andrew Ng**

Prompt engineering types: Reasoning & Acting (ReAct)

- ReAct is technique which enable LLMs to do reasoning and take task specific actions.
- Combines chain of thought reasoning with action planning.



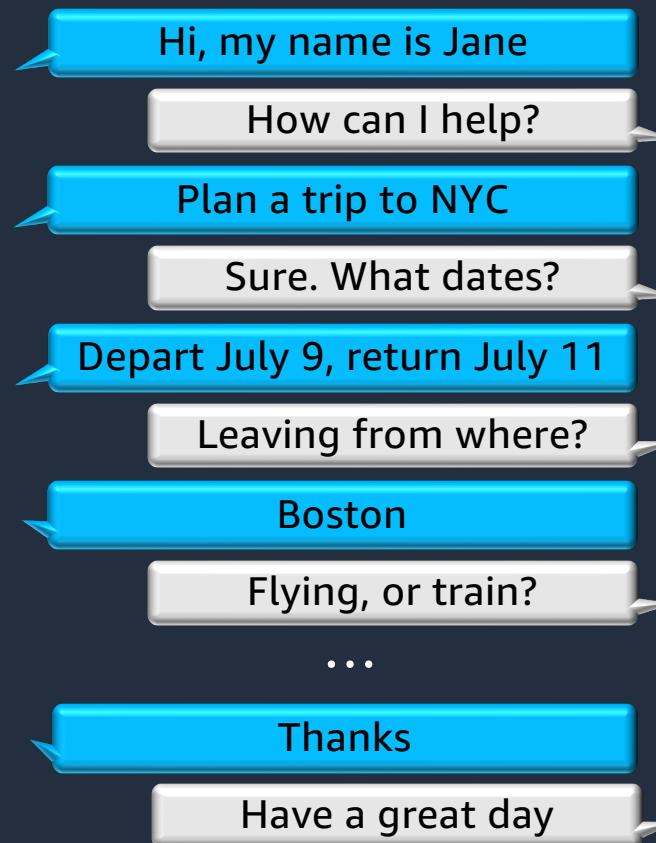
Agent orchestration



Agent breaks task into subtasks, determines the right sequence, and executes actions and knowledge searches on the fly

Agents use turn-by-turn dialog within a session to provide a natural experience

Jane books a trip using an agent



Travel
advisor
agent

Jane uses the agent
for ~10 minutes,
~20 turns



Math calculations using code interpretation, no files



Mortgage agent



What is my estimated monthly payment for a \$500K 30-year fixed loan?

The current rate for that loan type is 8%. Your payment would be \$3,668.82.

Agent orchestration

1. Find the current interest rate in knowledge base
2. Generate code for calculating payments
3. Execute code to estimate payment amount
4. Use results in final response

Rates lookup

Rates knowledge base

Generated calculation



Code interpreter

How Computer Use works



1. Provide Claude with computer use tools and a user prompt

- Add Anthropic-defined computer use tools to your API request.

```
◦ { "type": "computer_20241022", "name": "computer" }  
◦ { "type": "text_editor_20241022", "name": "str_replace_editor" }  
◦ { "type": "bash_20241022", "name": "bash" }
```

- Include a user prompt that require these tools, e.g., “Go to Amazon.com and add diapers to my cart”



2. Claude decides to use a tool

- Claude loads the stored computer use tool definitions and assesses if any tools can help with the user's query.
- If yes, Claude constructs a properly formatted tool use request.
- The API response has a **stop_reason** of **tool_use**, signaling Claude's intent.



How Computer Use works



3. Extract tool input, evaluate the tool on a computer, and return results

- On the client side, extract the tool name and input from Claude's request



Use a dedicated virtual machine or container with minimal privileges to prevent direct system attacks or accidents



4. Claude continues calling computer use tools until it's completed the task

- Claude analyzes the tool results to determine if more tool use is needed or the task has been completed.

Summary

- Prompt Types
- Retrieval Augmented Generation
- Fine Tuning
- Agents
- Computer Use





Thank you!

Praveen Jayakumar
Head of AI/ML Solution Architecture
AWS India