



Asia's Largest

AI & Cloud Conference 2024

15 - 16, November 2024

Chennai Trade Center, Chennai



Topic

Unraveling The Microservices Maze: Advanced Distributed Tracing In Kubernetes



Who I am

Vijayperumal

Technical Architect @



Episode #01 - Microservices

1) API Gateway

The gateway provides a unified entry point for client applications. It handles routing, filtering, and load balancing.

2) Service Registry

The service registry contains the details of all the services. The gateway discovers the service using the registry. For example, Consul, Eureka, Zookeeper, etc.

3) Service Layer

Each microservices serves a specific business function and can run on multiple instances. These services can be built using frameworks like Spring Boot, NestJS, etc.

4) Authorization Server

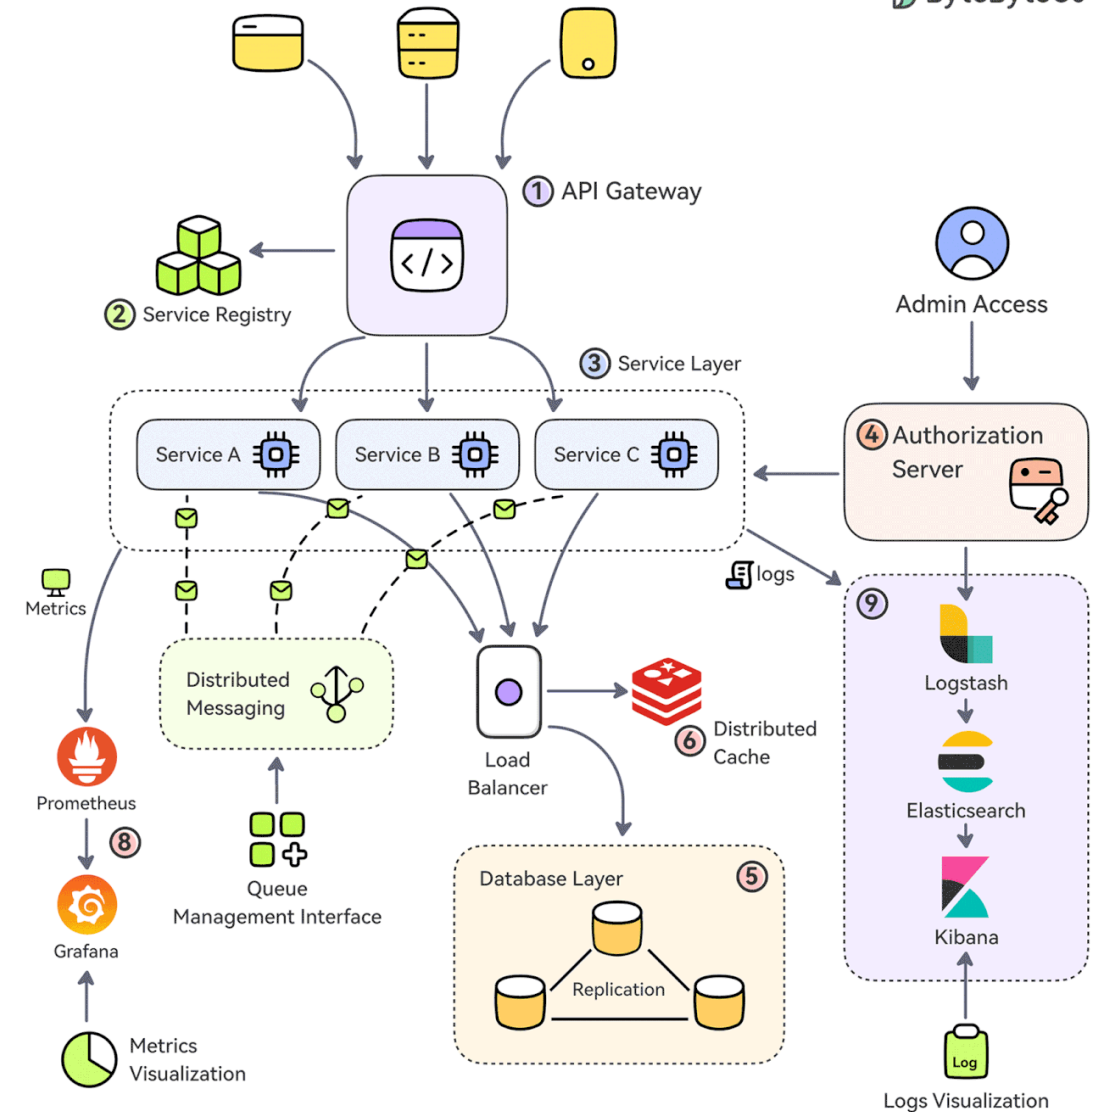
Used to secure the microservices and manage identity and access control. Tools like Keycloak, Azure AD, and Okta can help over here.

5) Data Storage

Databases like PostgreSQL and MySQL can store application data generated by the services.

9 Essential Components of Production Microservice App

ByteByteGo



5) Distributed Caching

Caching is a great approach for boosting the application performance. Options include caching solutions like Redis, Couchbase, Memcached, etc.

6) Async Microservices Communication

Use platforms such as Kafka and RabbitMQ to support async communication between microservices.

7) Metrics Visualization

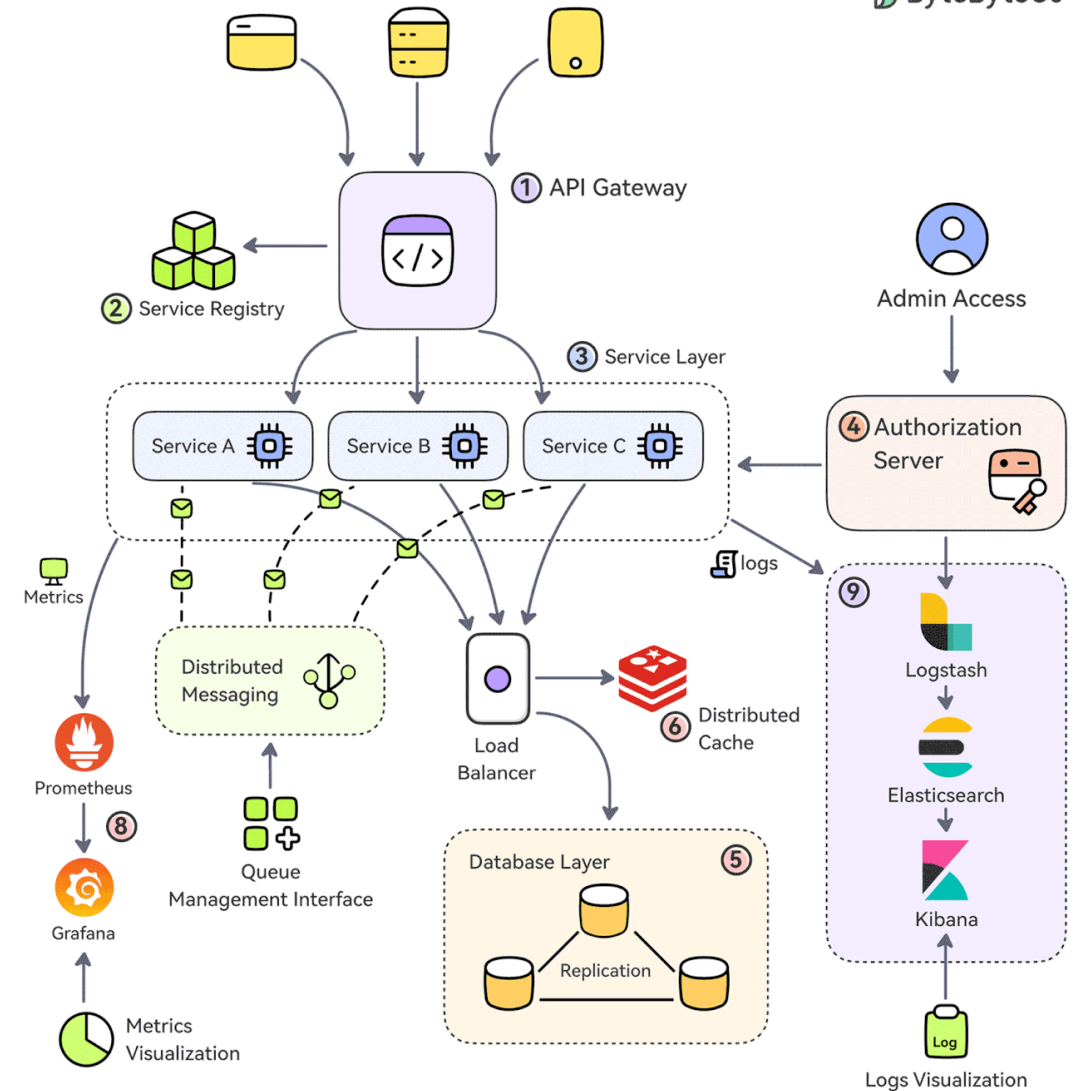
Microservices can be configured to publish metrics to Prometheus and tools like Grafana can help visualize the metrics.

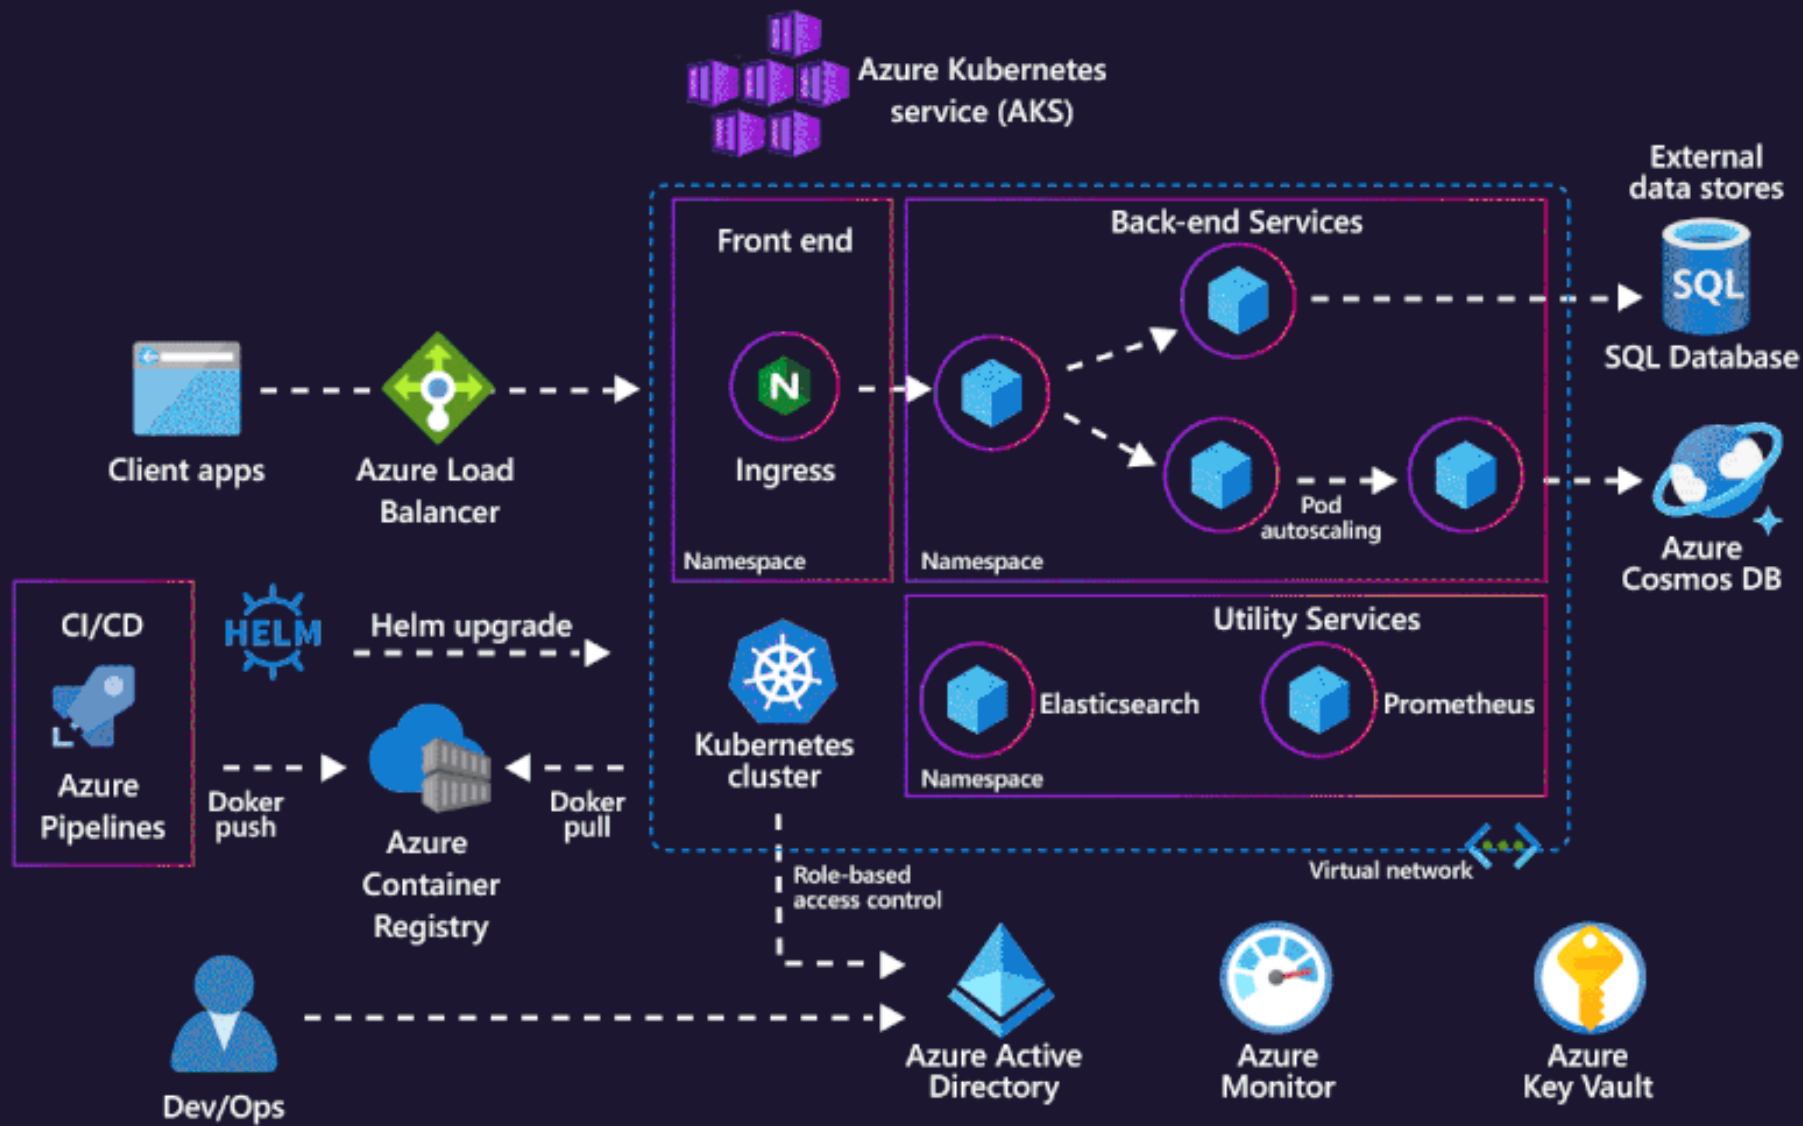
8) Log Aggregation and Visualization

Logs generated by the services are aggregated using Logstash, stored in Elasticsearch, and visualized with Kibana.

9 Essential Components of Production Microservice App

ByteByteGo





Microservice Challenges

Using Kubernetes



Complexity of Orchestration

Managing multiple services, pods, and replicas.



Observability

Difficulties in monitoring distributed systems.



Networking and Service Discovery

Inter-service communication challenges



Resource Management

Efficiently allocating CPU and memory.



Security

Securing microservices and APIs in a dynamic environment



Microservice Challenges

Using Kubernetes



Complexity of Orchestration

Managing multiple services, pods, and replicas.



Observability

Difficulties in monitoring distributed systems.



Networking and Service Discovery

Inter-service communication challenges



Resource Management

Efficiently allocating CPU and memory.



Security

Securing microservices and APIs in a dynamic environment



Microservice Challenges

Using Kubernetes



Complexity of Orchestration

Managing multiple services, pods, and replicas.



Observability

Difficulties in monitoring distributed systems.



Networking and Service Discovery

Inter-service communication challenges



Resource Management

Efficiently allocating CPU and memory.



Security

Securing microservices and APIs in a dynamic environment



Microservice Challenges

Using Kubernetes



Complexity of Orchestration

Managing multiple services, pods, and replicas.



Observability

Difficulties in monitoring distributed systems.



Networking and Service Discovery

Inter-service communication challenges



Resource Management

Efficiently allocating CPU and memory.



Security

Securing microservices and APIs in a dynamic environment



Microservice Challenges

Using Kubernetes



Complexity of Orchestration

Managing multiple services, pods, and replicas.



Observability

Difficulties in monitoring distributed systems.



Networking and Service Discovery

Inter-service communication challenges



Resource Management

Efficiently allocating CPU and memory.



Security

Securing microservices and APIs in a dynamic environment






Episode #02 - Observability

Observability


Observability is the ability to understand the internal state of a system by examining its outputs.

In the context of software, this means being able to understand the internal state of a system by examining its telemetry data, which includes traces, metrics, and logs.




BlackBox

We can't understand a complex system if it's a black box.



WhiteBox

The only way to light up those black boxes is with high-quality telemetry: distributed traces, metrics, logs, and more.



Outputs -> Telemetry Data

To make a system observable, it must be instrumented. That is, the code must emit signals or telemetry data such as: traces, metrics, or logs (MELT for short). The instrumented data must then be sent to an **observability backend**.

The three core elements of data observability are:



Logs



Metrics



Tracing


Observability Benefits



Higher
Visibility



Speed up
Troubleshooting



Team
Productivity



Finding out
unknown
issues




Help to
improve end
user experience



Help to
reduce costs

People benefiting from observability

Many different teams can use observability to understand the behavior of complex digital systems and turn data into tailored insights.



Monitoring X Observability

Based on predefined sets of metrics or logics.

What? + When?

E.g.

- What to expect from the software?
- When did the software present an unexpected behaviour?

Monitoring presents the software behaviour

Based on exploring properties and patterns not defined in advance.

Why?

E.g.

- Why is the software behaving unexpectedly?

Observability explains the software behavior



Episode #03 - Open Telemetry

What is Open Telemetry

OpenTelemetry, also known as OTel, is a vendor-neutral open source Observability framework

OpenTelemetry is focused on the generation, collection, management, and export of telemetry data such as traces, metrics, and logs

Two key principles:

1. You own the data that you generate.
There's no vendor lock-in
2. You only have to learn a single set of
APIs and conventions

Sponsored by the
Support for different programming languages



What Type of Telemetry Data Does OpenTelemetry Handle?

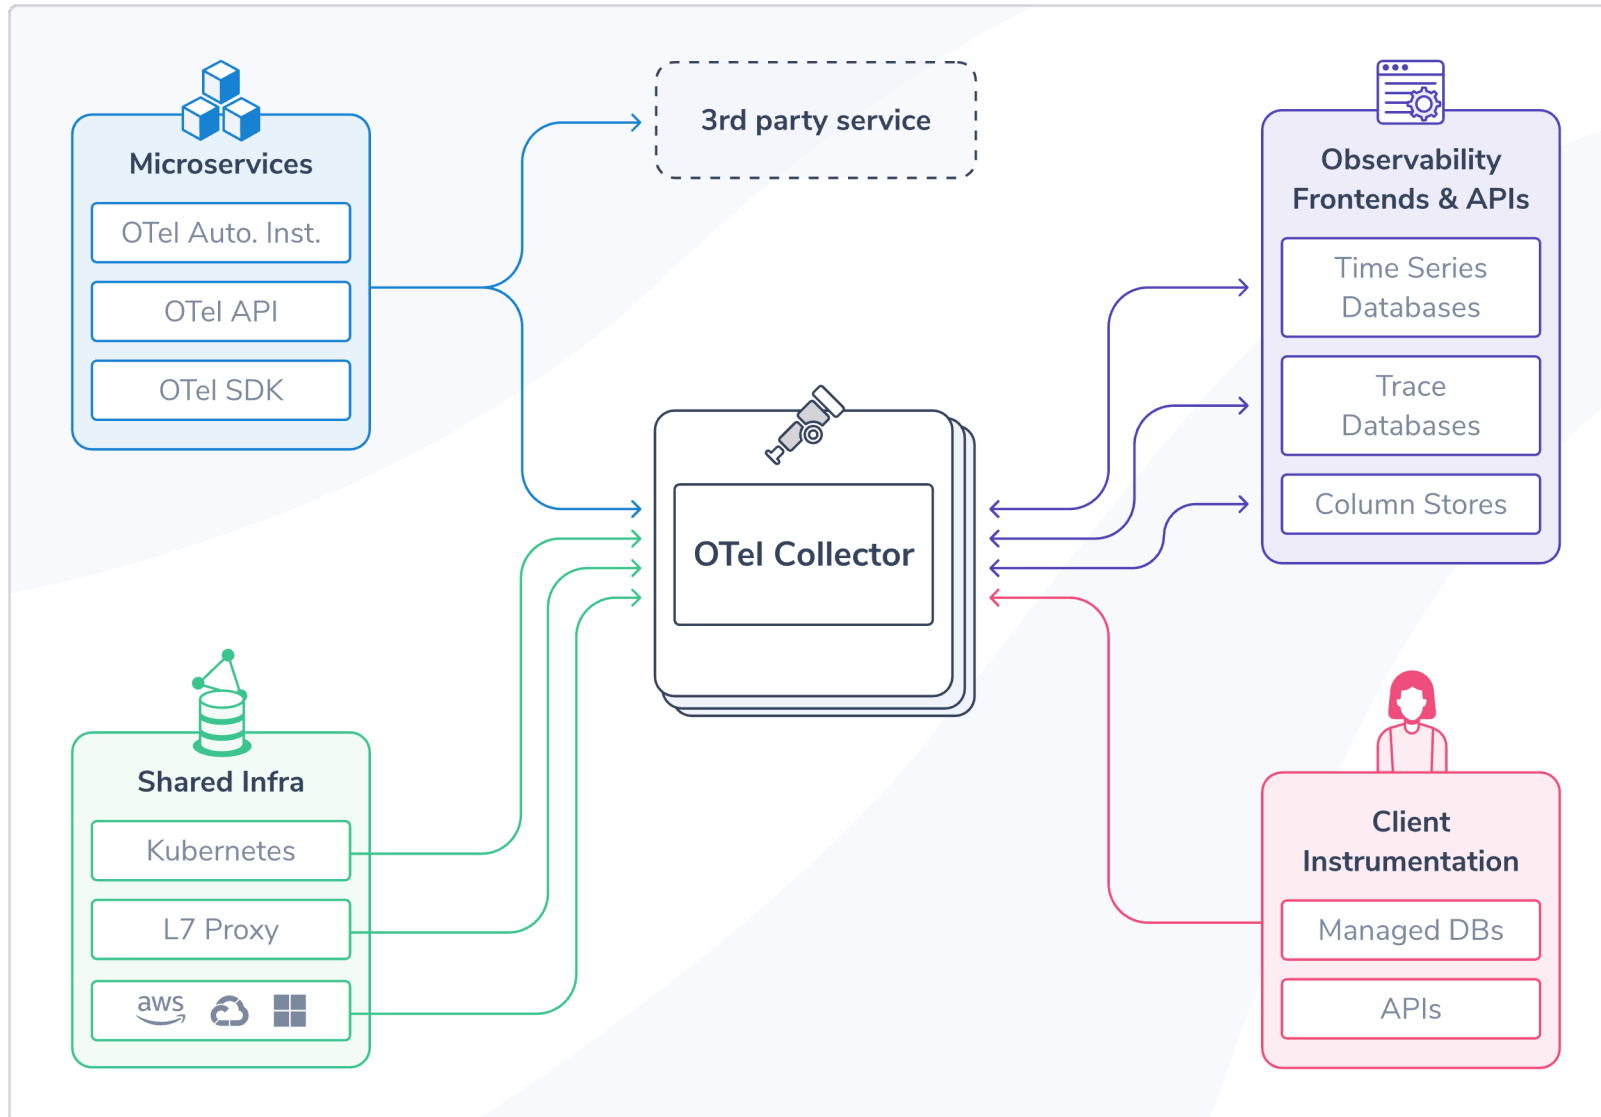
Open Telemetry handles three primary types of telemetry data:

Traces allow developers to track the journey of a request through various services, helping to identify bottlenecks and understand the flow of requests within a system.

Metrics provide quantitative information about the operation of applications and infrastructure, such as response times, memory usage, and request counts, enabling performance monitoring and trend analysis.


Logs offer qualitative insights through event records, detailing what happened in the system at a specific point in time. Together, these data types offer a holistic view of system performance and behavior, aiding in debugging, monitoring, and optimizing applications.

OpenTelemetry Architecture and Components



Integrations

Standards




Tracers




Service Meshes / Proxies



Languages




Containers, Platforms and Clouds




Data Streaming and Storage



Deployment Automation (CI/CD)



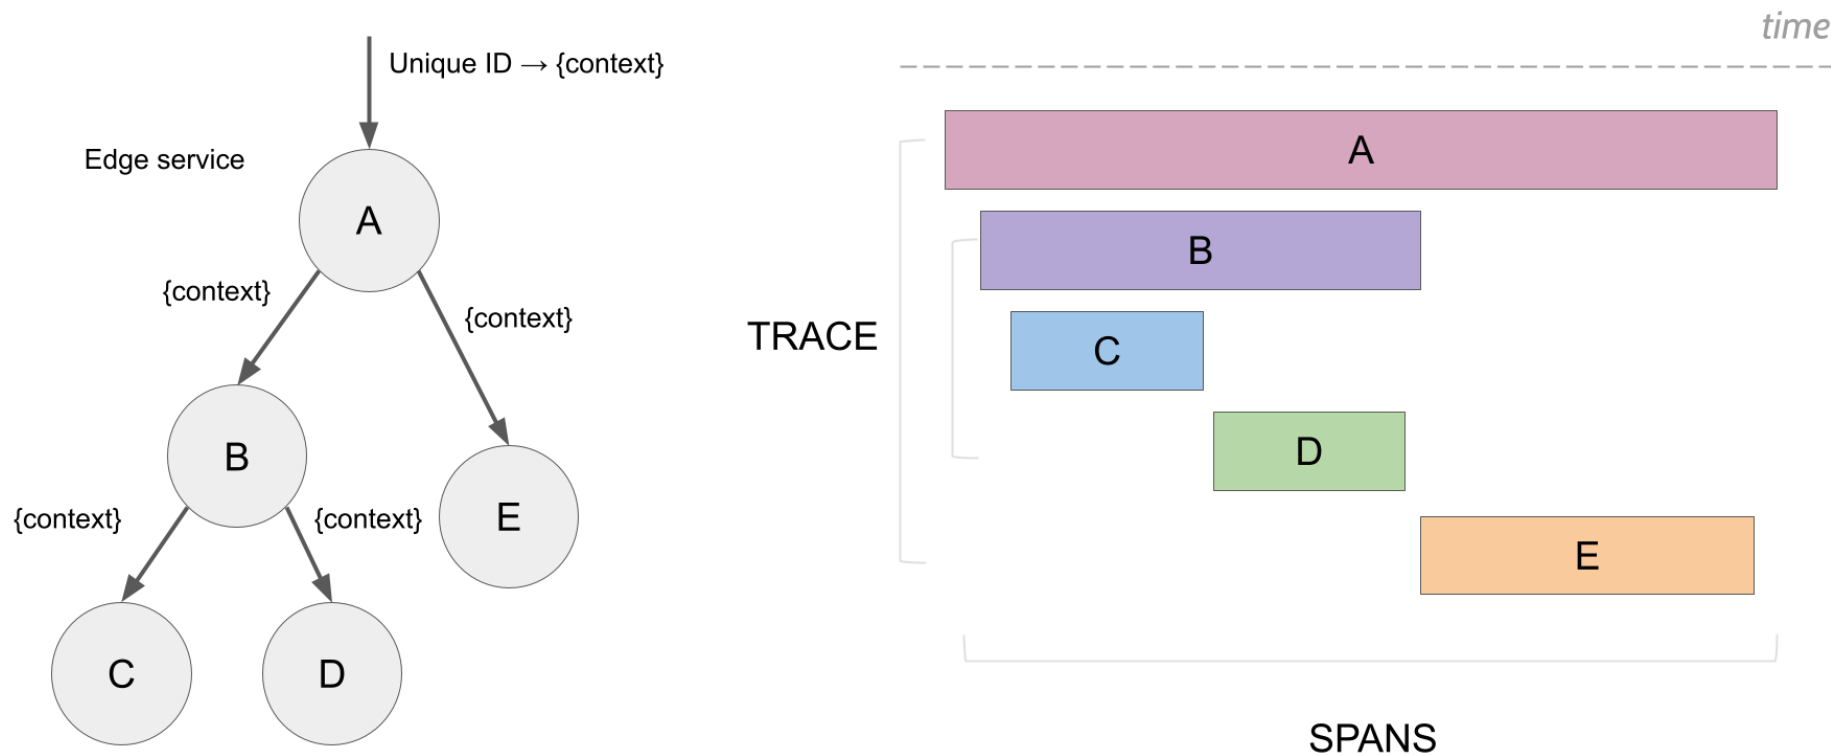
Alerting and Tools



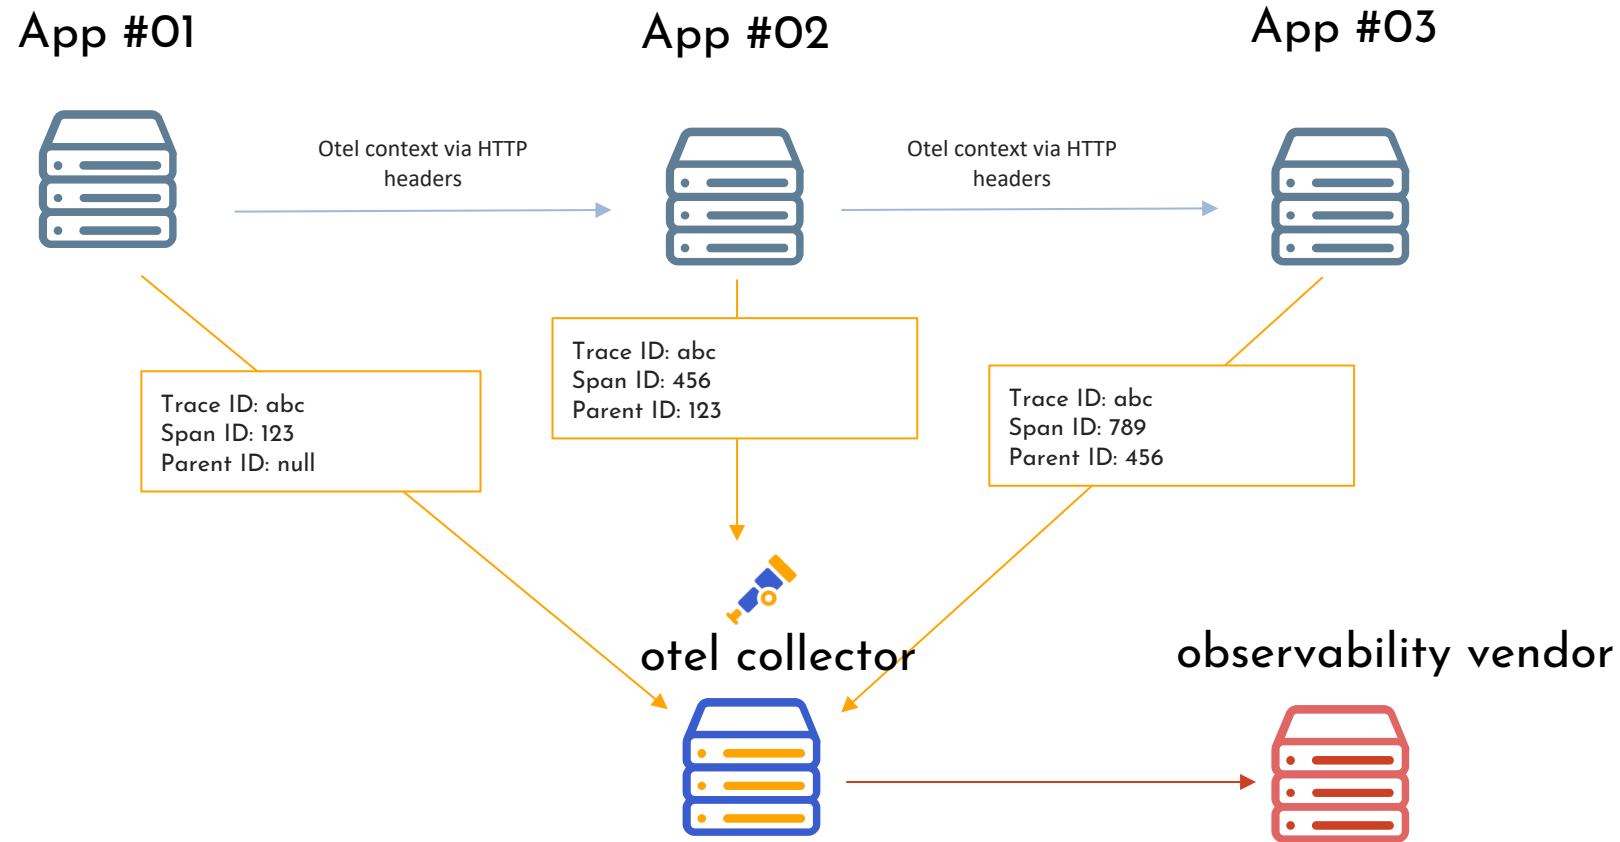
Terminology

A **span** represents a logical unit of work in Jaeger that has an operation name, the start time of the operation, and the duration. Spans may be nested and ordered to model causal relationships.

A **trace** is a data/execution path through the system, and can be thought of as a directed acyclic graph of spans.



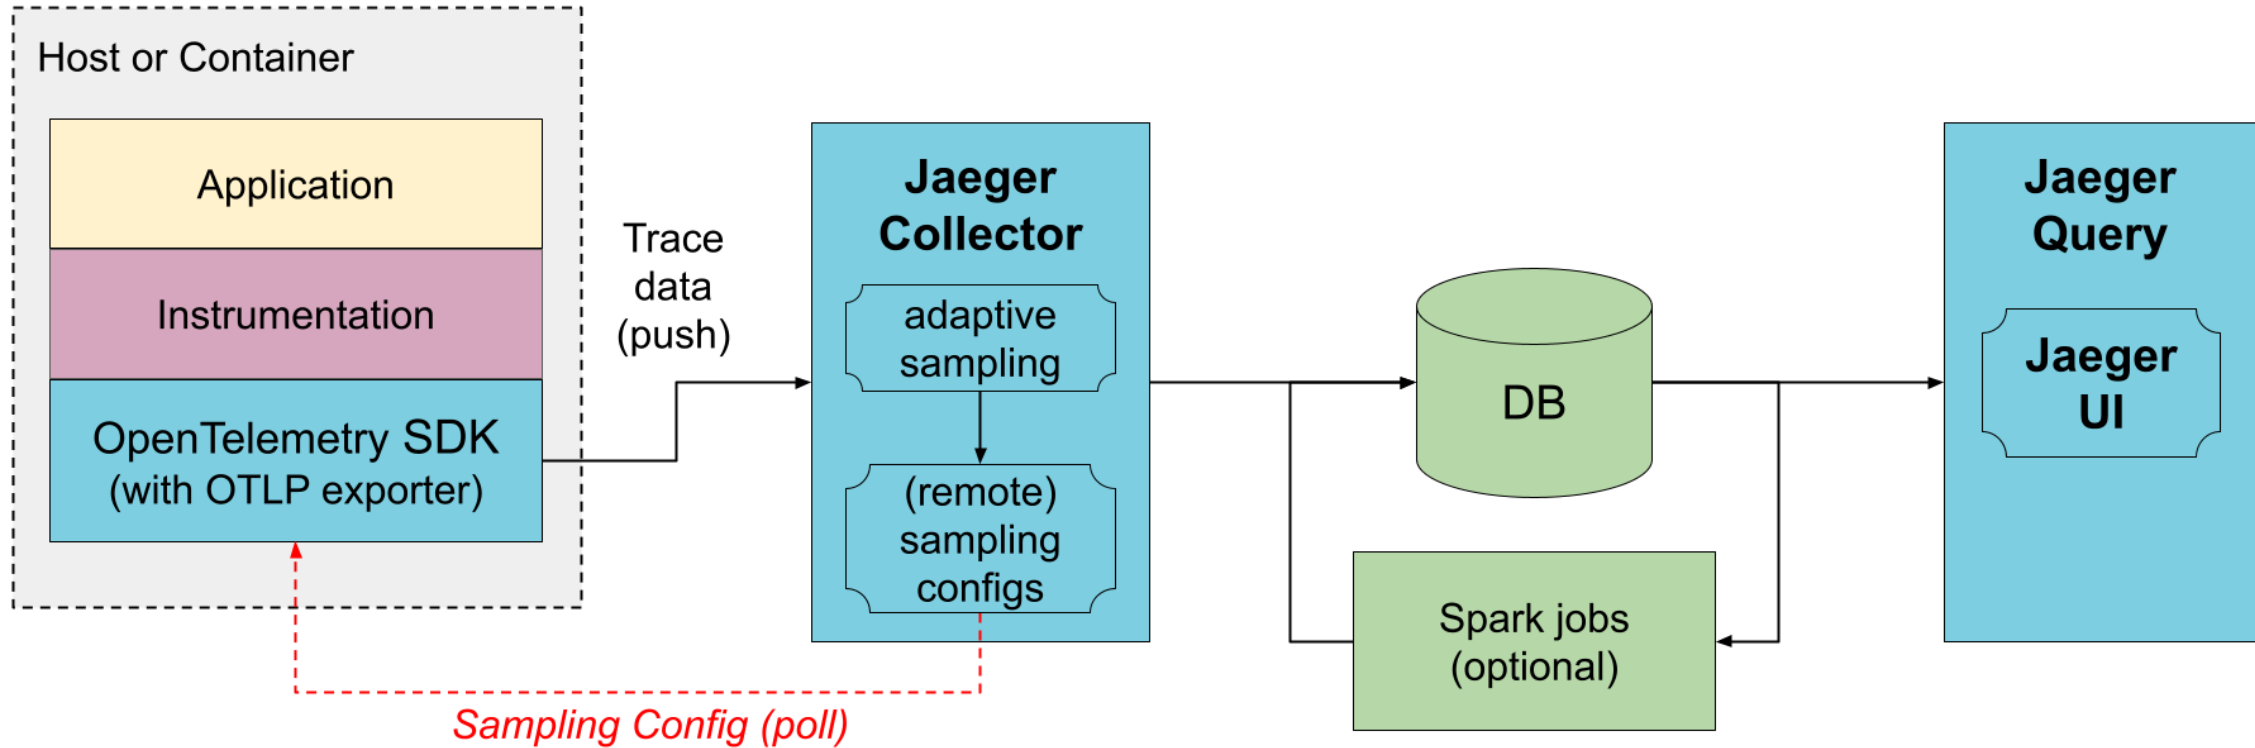
Distributed Tracing





Episode #04 - Jaeger UI

Architecture






Episode #07 - Best Practices

Best Practices for Distributed Tracing in Production

- ✓ **Optimize Sampling Strategies**
- ✓ **Tag Meaningfully and Consistently**
- ✓ **Reduce Tracing Overhead**
- ✓ **Monitor Trace Health and Coverage**
- ✓ **Leverage Dashboards and Alerts**
- ✓ **Integrate Traces with Logs and Metrics**
- ✓ **Train and Align Teams**
- ✓ **Scale Efficiently**
- ✓ **Focus on Actionable Insights**
- ✓ **Choose the Right Tools and Ecosystem**





The barriers to adopting Distributed Tracing – and how AI is going to remove them.

The Configuration Issue

The Sampling Issue

The Trace Enrichment Issue

What you get

You get a full eight hours of sleep.
Your company closes the deal.
Everybody is happy.



Copilot LLM Prompt

Show me requests that had above average latency

Which requests threw errors or exceptions?

Did any of those requests have empty or above average payloads?



Episode #06 - The Review



The End - *See you all TechXConf 2024*

